

Reversing SR-IOV For Fun and Profit

Adir Abraham

5-8 September, 2018

r2con




Legal Notices and disclaimers

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer or retailer or learn more at [intel.com].

Intel disclaims all express and implied warranties, including without limitation, the implied warranties of merchantability, fitness for a particular purpose, and non-infringement, as well as any warranty arising from course of performance, course of dealing, or usage in trade.

Intel and the Intel logo, are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.



\$ whoami

- ▶ Name: Adir Abraham
- ▶ Title: SW/FW Security Researcher & Reverse Engineer @ Intel
- ▶ Graduated: BSc Computer Science Edu.; BA Economics - Technion IIT
- ▶ Certified: CISSP, CySA+, CCSK

- ▶ Twitter: [@adirab](https://twitter.com/adirab)
- ▶ LinkedIn: <https://linkedin.com/in/adirab>

Agenda

▶ Overview

- ▶ Introduction
 - ▶ PCIe Devices
 - ▶ I/O Virtualization

▶ PCIe Virtualization

- ▶ Motivation
- ▶ Directed I/O
- ▶ PCIe Architecture

▶ SR-IOV

- ▶ Architecture Supporting SR-IOV Capability
- ▶ ARI - Alternative Routing ID Interpretation
- ▶ ATS - Address Translation Service

▶ Demo

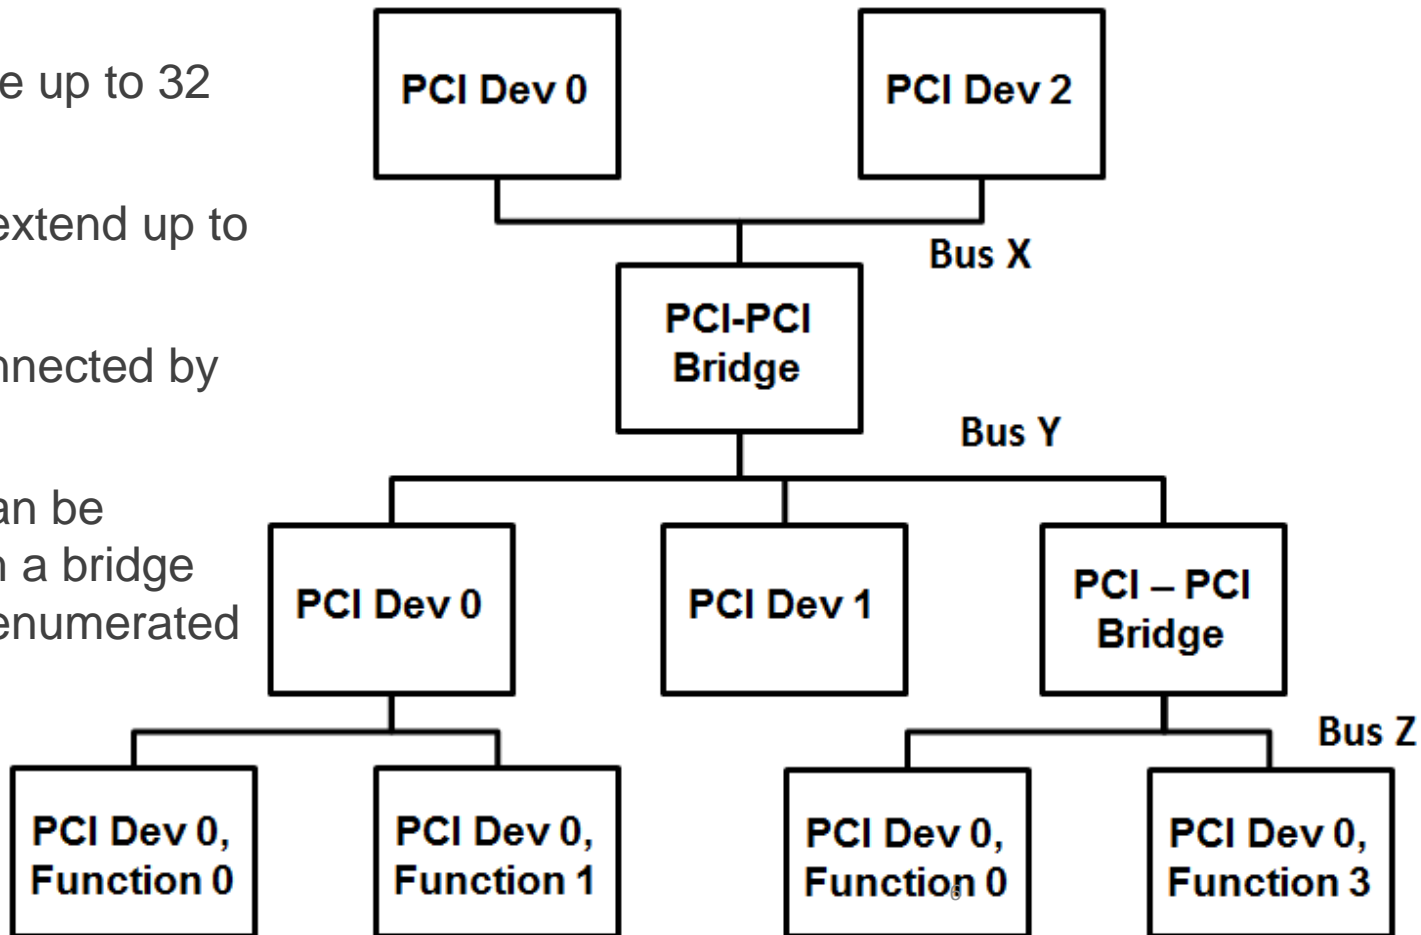
- ▶ Finding PCIe features using radare2
- ▶ Finding SR-IOV capabilities using radare2

PCI - Peripheral Component Interconnect

- ▶ It's a bus protocol developed by Intel around 1993
- ▶ Purpose is to attach/interconnect local hardware devices to a computer
- ▶ PCI is integrated into the chipset, forming a “backbone”
 - ▶ Holds true for both Intel and AMD-based systems
 - ▶ Logically speaking, the Chipset is a PCI System
- ▶ 32-bit bus with multiplexed address and data lines
 - ▶ Supports 64-bit by performing two 32-bit reads
- ▶ PCI component interface is processor independent
 - ▶ The CPU/BIOS reads and writes to the configuration space to configure much of the system
- ▶ Intel's MCH/ICH chipsets implement PCI Local Bus Protocol 2.3
- ▶ PCH chipsets implement PCI Express protocol (v. 2.0)
 - ▶ Still supports PCI 2.3
- ▶ PCI standards are currently maintained and defined by the PCI /SIG:
 - ▶ <http://www.pcisig.com/specifications/>

Generic PCI Topology: Buses, Devices, and Functions

- ▶ Up to 256 Buses
- ▶ Each Bus can have up to 32 devices attached
- ▶ Each Device can extend up to 8 Functions
- ▶ Buses are interconnected by Bridges
- ▶ Multiple bridges can be connected through a bridge interface and are enumerated




PCI Address Spaces

- ▶ PCI implements three address spaces:
- ▶ 1) PCI Configuration Space (up to 256 Bytes)
 - ▶ Required/standard. Defined in the specifications. Every PCI device has a configuration space.
- ▶ 2) PCI Memory-mapped space
 - ▶ Optional. Dependent on whether the device manufacturer needs to map system memory to the PCI device
- ▶ 3) PCI I/O-mapped space
 - ▶ Optional. Same as PCI Memory Space

PCI Express (PCIe)

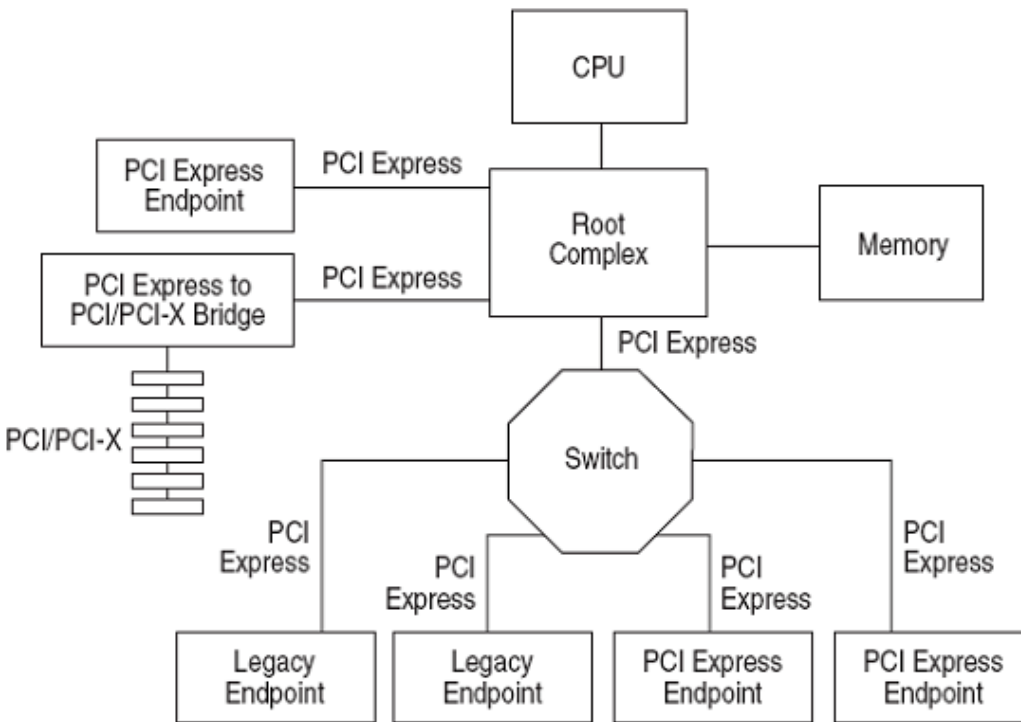
- ▶ Peripheral Component Interconnect Express
- ▶ Developed around 2004
- ▶ Packet-based transaction protocol
- ▶ Very different from PCI at the hardware level
- ▶ For software configuration purposes, it is mostly the same
 - ▶ Adds an extended configuration space of 4KB
- ▶ Provides backwards compatibility for Compatible PCI
- ▶ Adds 4KB of PCI Express Extended Capabilities registers
 - ▶ Located in either the Configuration space starting at offset 256 (immediately following the Compatible PCI configuration space)
 - ▶ Or located at a MMIO location specified in the Root Complex Register Block (RCRB)



PCI Express (PCIe) Address Spaces

- ▶ PCIe implements four address spaces:
- ▶ 1) PCIe Configuration Space (up to 4KBytes)
 - ▶ Required/standard. Defined in the specifications. Every PCIe device has its configuration space mapped to memory.
 - ▶ Also provides the first 256 bytes of compatible PCI (memory-mapped and via port IO for backwards compatibility)
- ▶ 2) PCIe Memory-mapped space
 - ▶ Optional. Dependent on whether the device manufacturer needs to map system memory to the PCI device
- ▶ 3) PCIe I/O-mapped space
 - ▶ Optional. Same as PCI Memory Space
- ▶ 4) PCIe Message Space
 - ▶ For low-level protocol messaging/interrupts. We don't get into this in this class

Generic PCIe Topology:




- ▶ The Root Complex connects the processor to the system memory and components
- ▶ Same number of devices supported as PCI
- ▶ Up to 256 PCIe buses
- ▶ Up to 32 PCIe devices
- ▶ Up to 8 Functions
- ▶ Each Function can implement up to 4 KB of configuration space

I/O Virtualization Concept


- ▶ Virtualizing the I/O path between a host and a PCIe device
- ▶ **Prevent** direct access to physical (memory) resources
- ▶ **Managing** and **separating** each device's I/O space using a VMM
- ▶ Can apply to anything that uses an adapter in a computer, such as:
 - ▶ Ethernet Network Interface Cards (NICs)
 - ▶ Disk Controllers (including RAID controllers)
 - ▶ Fibre Channel Host Bus Adapters (HBAs)
 - ▶ Graphics/Video cards or co-processors
 - ▶ SSDs mounted on internal cards

I/O Virtualization Implementations


- ▶ A - Software only (emulation)
- ▶ B - Directed I/O (enhance performance)
- ▶ C - Directed I/O and Device Sharing (resource saving)



A - Software only



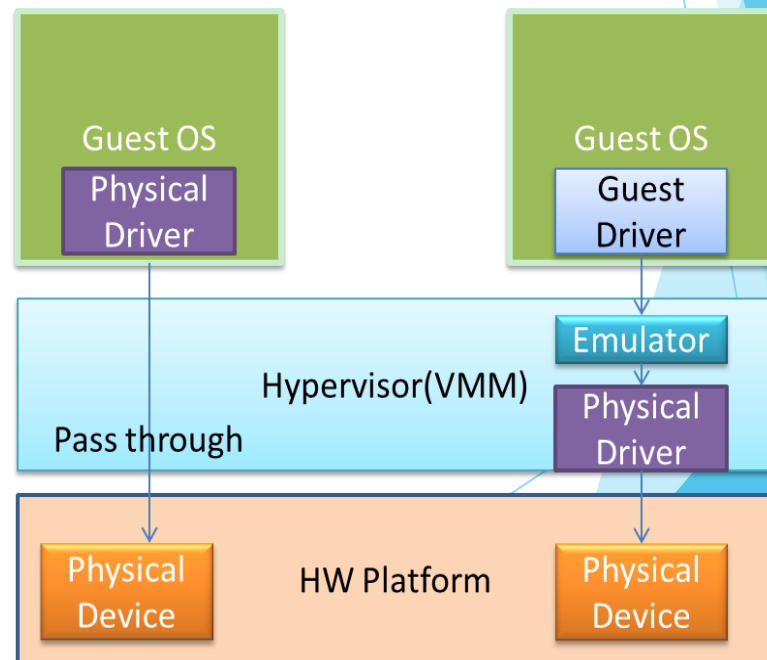
B - Directed I/O



C - Directed I/O & Device Sharing

Directed I/O - Passthrough

- ▶ Software-based sharing adds overhead to each I/O due to emulation layer
 - ▶ This indirection has the additional affect of eliminating the use of hardware acceleration that may be available in the physical device.
- ▶ Directed I/O has added enhancements to facilitate memory translation and ensure protection of memory that enables a device to directly DMA to/from host memory.
 - ▶ Bypass the VMM's I/O emulation layer
 - ▶ Throughput improvement for the VMs



Drawbacks to Directed I/O

- ▶ One concern with direct assignment is that it has limited scalability
- ▶ Assignment is per physically-isolated functionality
 - ▶ A **physical** device can only be assigned to **one** VM.
 - ▶ For example, a dual port NIC allows for direct assignment to two VMs. (one port per VM)
 - ▶ Consider for a moment a fairly substantial server of the very near future
 - ▶ 4 physical CPU's
 - ▶ 12 cores per CPU
 - ▶ If we use the rule that one VM per core, it would need 48 physical ports.

Terminology relating to Directed I/O

Acronym	Expansion	Defined By	What is it?
I/O MMU	I/O Memory Management Unit	Common parlance	Translation mechanism in the system memory controller (North Bridge) that allows a device or set of devices to use translated addresses when accessing main memory. In many cases, it also translates interrupts coming from the devices as messages.
ATPT	Address Translation and Protection Table	PCI SIG	I/O MMU
VT-d, VT-d2	Virtualization Technology for Directed I/O	Intel	I/O MMU
DMAR	DMA Remapping	Intel, Microsoft	I/O MMU
IOMMU	I/O Memory Management Unit	AMD	I/O MMU

Is IOMMU (vt-d) enabled on your system?

```
$ sudo dmesg | grep -e dmar -e IOMMU
[ 0.000000] DMAR: IOMMU enabled
[ 0.004000] DMAR: dmar0: reg_base_addr ZZZZZZZZ ver 1:0 cap XXXXXXXXXXXX ecap wwwwww
[ 0.004000] DMAR-IR: IOAPIC id 128 under DRHD base 0xYYYYYYYY IOMMU 0
[ 1.192371] DMAR: dmar0: Using Queued invalidation
[ 1.526166] AMD IOMMUv2 driver by Joerg Roedel <jroedel@suse.de>
[ 1.526167] AMD IOMMUv2 functionality not available on this system
```

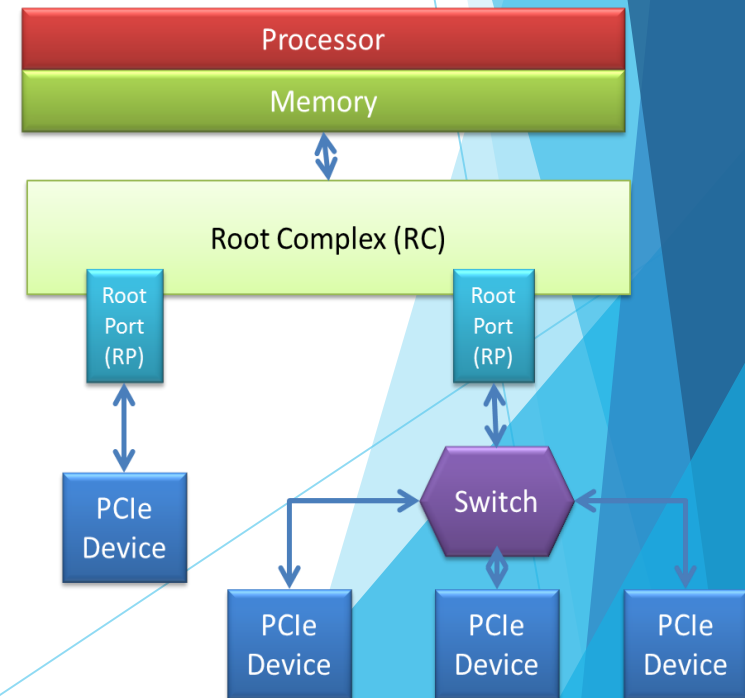
IOMMUv2 includes:

- (a) I/O device assignment
- (b) DMA remapping
- (c) interrupt remapping
- (d) reliability features

PCIe components

▶ Root Complex

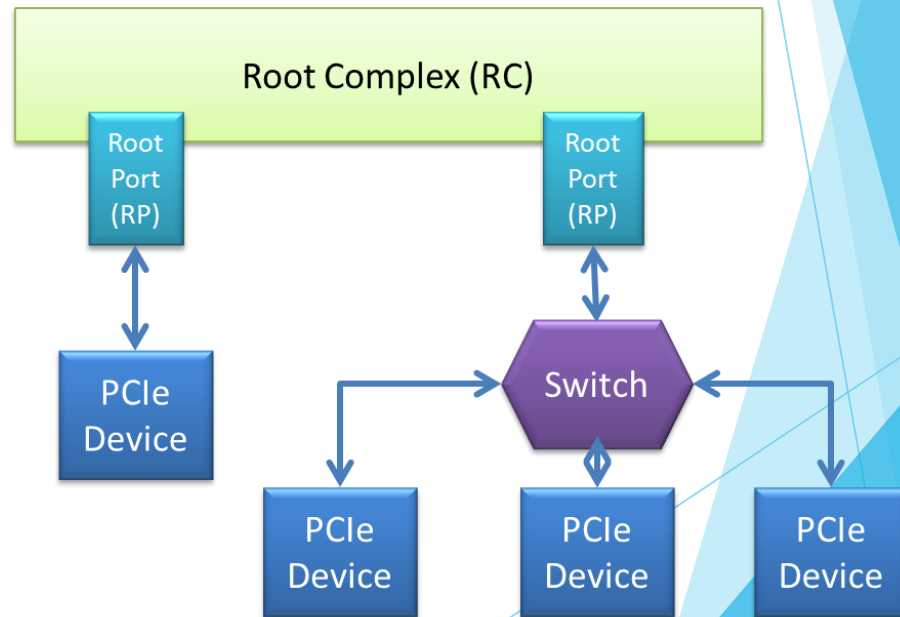
- ▶ A root complex connects the processor and memory subsystem to the PCIe switch fabric composed of one or more switch devices
- ▶ Similar to a **host bridge** in a PCI system
 - ▶ Generate transaction requests on behalf of the processor, which is interconnected through a local bus.
 - ▶ May contain more than one PCIe port and multiple switch devices.



PCIe components

▶ Root Port (RP)


- ▶ The portion of the motherboard that contains the host bridge. The host bridge allows the PCIe ports to talk to the rest of the computer



PCIe Device

- ▶ Unique PCI Function Address
 - ▶ Bus / Dev / **Function**
 - ▶ Command, lspci -v, can get PCI device information on linux

```
01:00.1 Ethernet controller: Intel Corporation I350 Gigabit Network Connection (rev 01)
Subsystem: Intel Corporation Ethernet Server Adapter I350-T2
Flags: fast devsel, IRQ 17
Memory at f7700000 (32-bit, non-prefetchable) [size=1M]
Memory at f7980000 (32-bit, non-prefetchable) [size=16K]
Capabilities: [40] Power Management version 3
Capabilities: [50] MSI: Enable- Count=1/1 Maskable+ 64bit+
Capabilities: [70] MSI-X: Enable- Count=10 Masked-
Capabilities: [a0] Express Endpoint, MSI 00
Capabilities: [100] Advanced Error Reporting
Capabilities: [140] Device Serial Number a0-36-9f-ff-ff-17-b1-6c
Capabilities: [150] Alternative Routing-ID Interpretation (ARI)
Capabilities: [160] Single Root I/O Virtualization (SR-IOV)
Capabilities: [1a0] Transaction Processing Hints
Capabilities: [1d0] Access Control Services
Kernel modules: igb
```



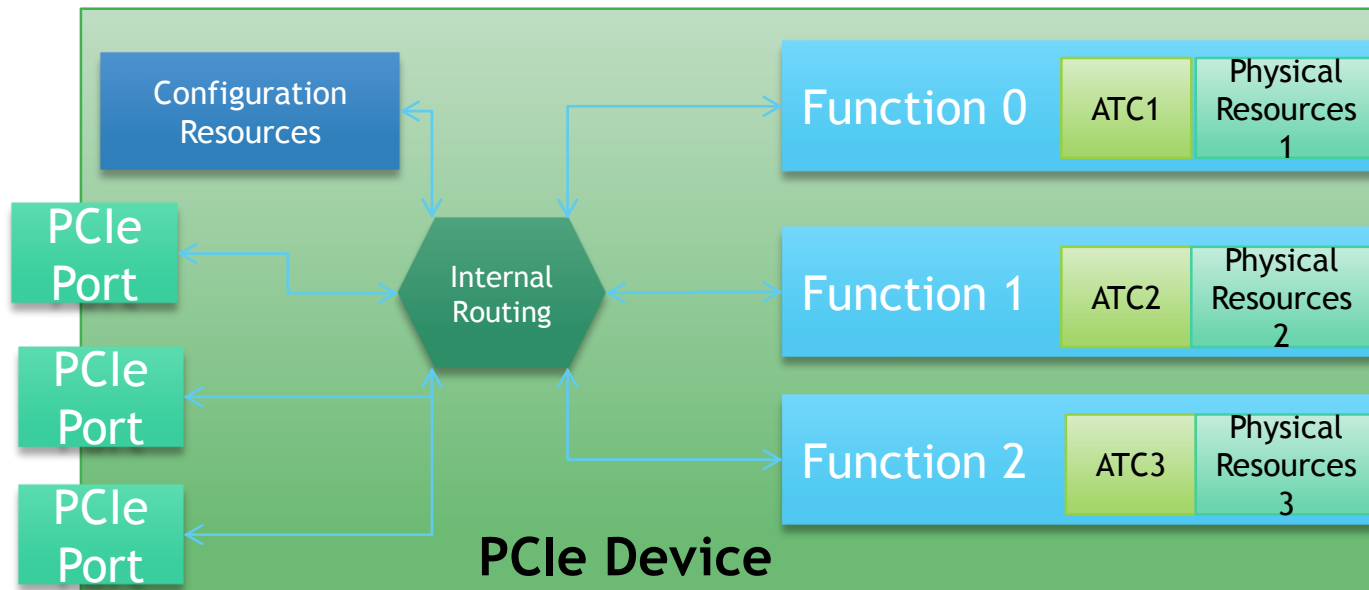
Device

Function1

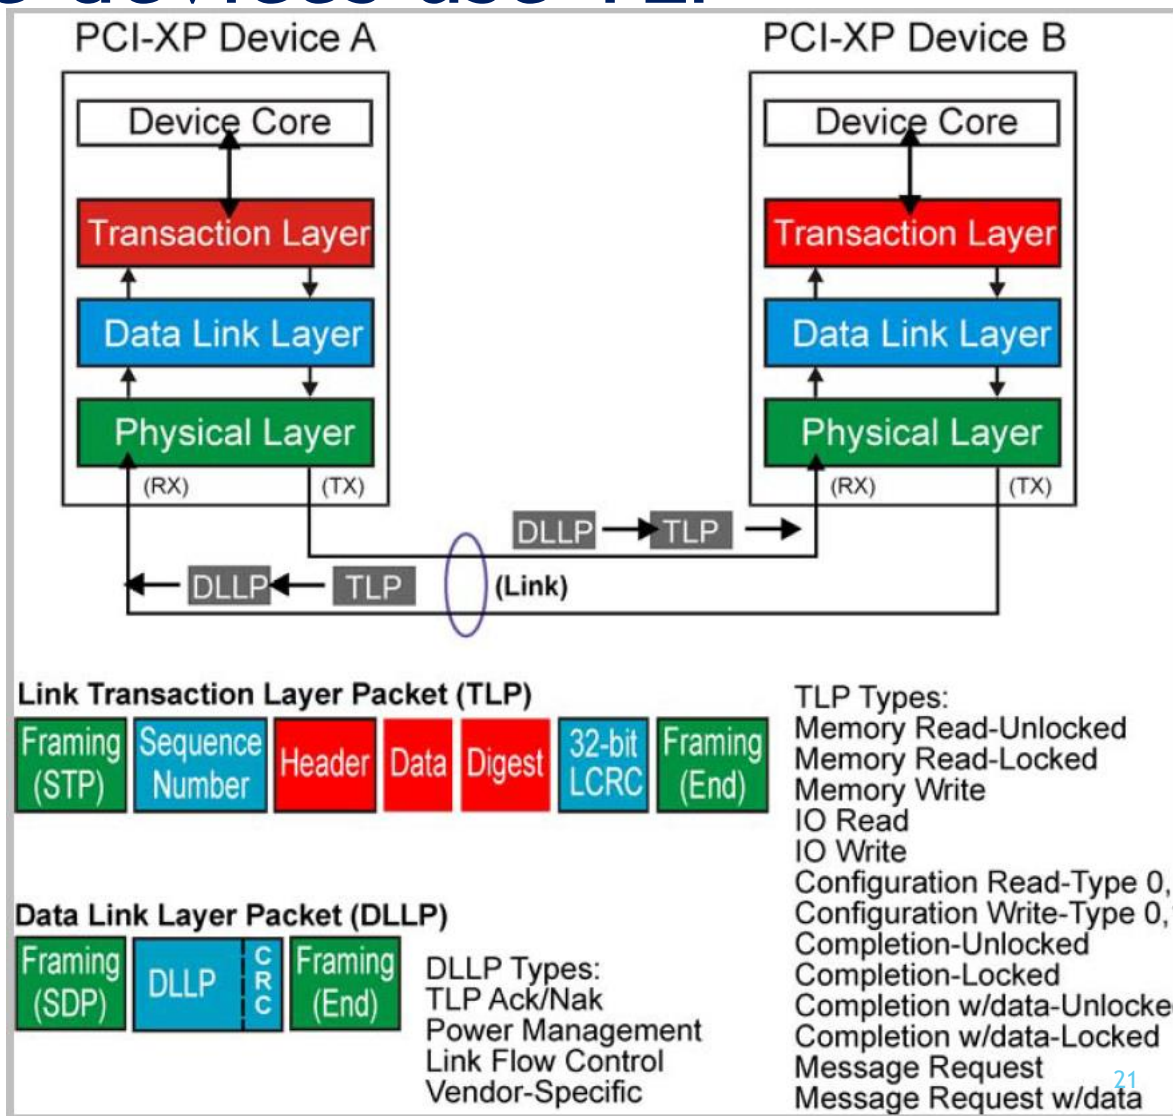
Function2

Example: Multi-Function Device

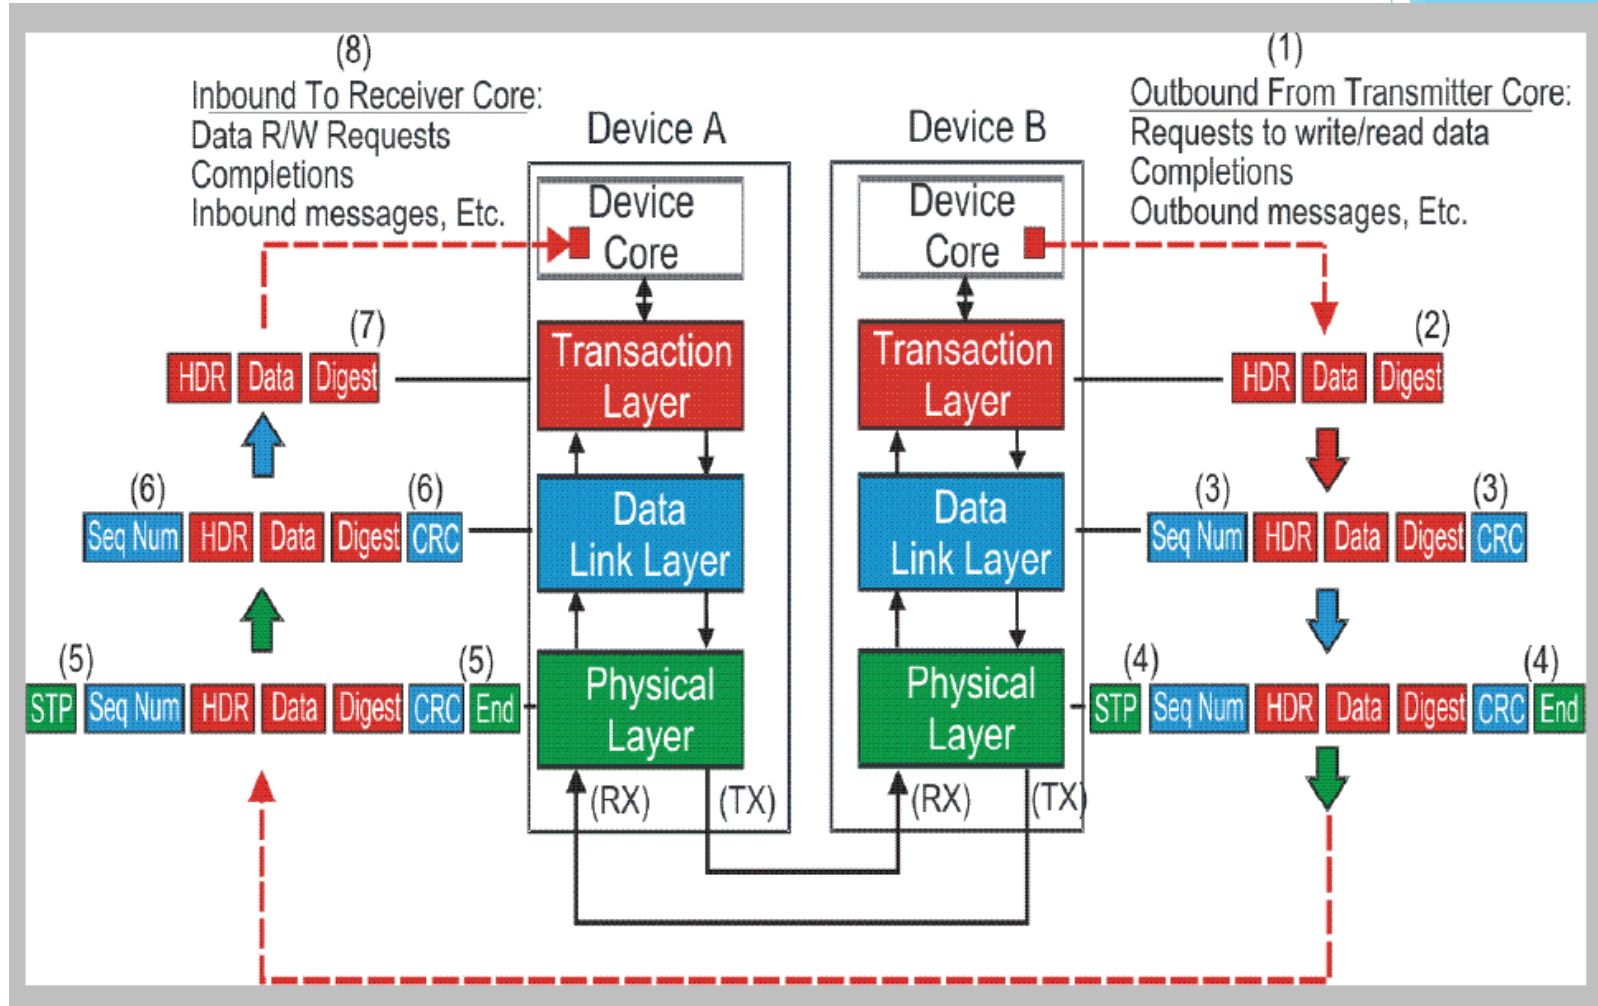
- ▶ The link and PCIe functionality shared by all functions is **managed through Function 0**
- ▶ All functions use a single Bus Number captured through the PCI enumeration process
- ▶ **Each function can be assigned to an SI**



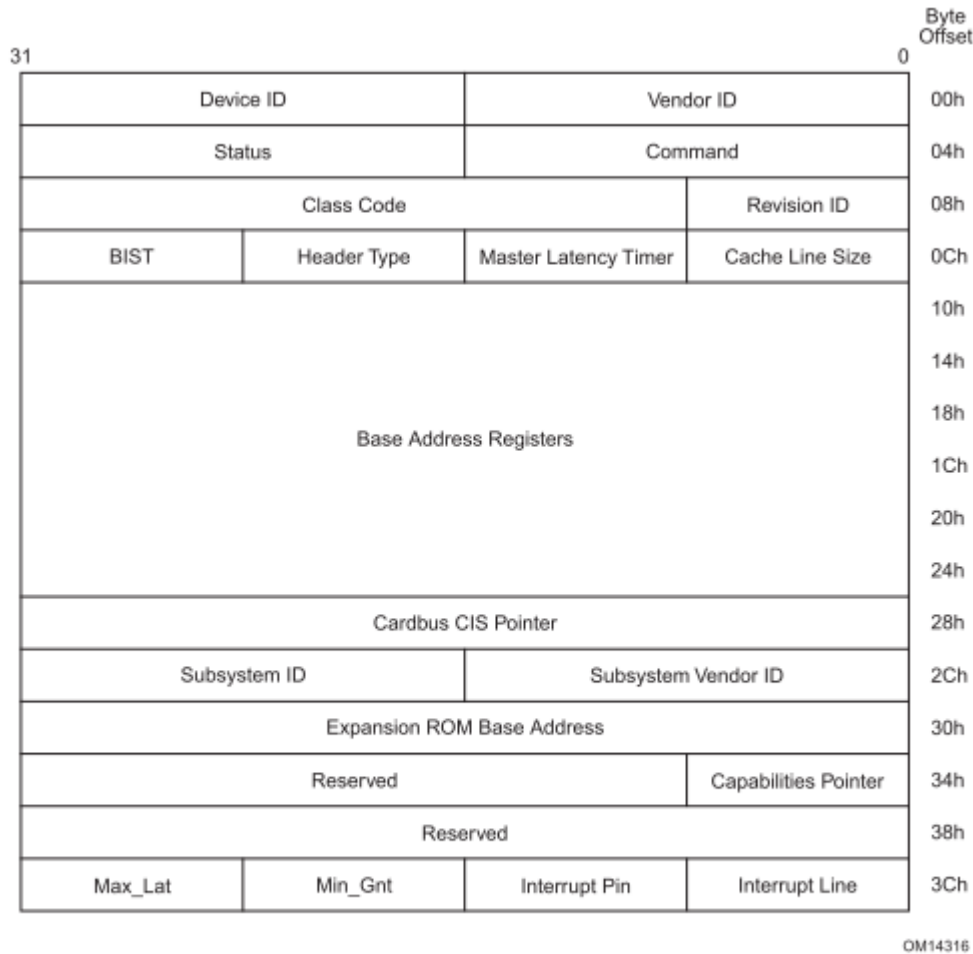
PCIe devices use TLP



TLP Assembly/Disassembly



Components in PCIe Device



Configuration Resources

► Configuration Space

- Devices will **allocate resource such as memory** and record the address into this configuration space

► Reference:

- PCI Local Bus Specification ver.2.3 Chap 6

Figure 3-2: Type 0 Configuration Space Header

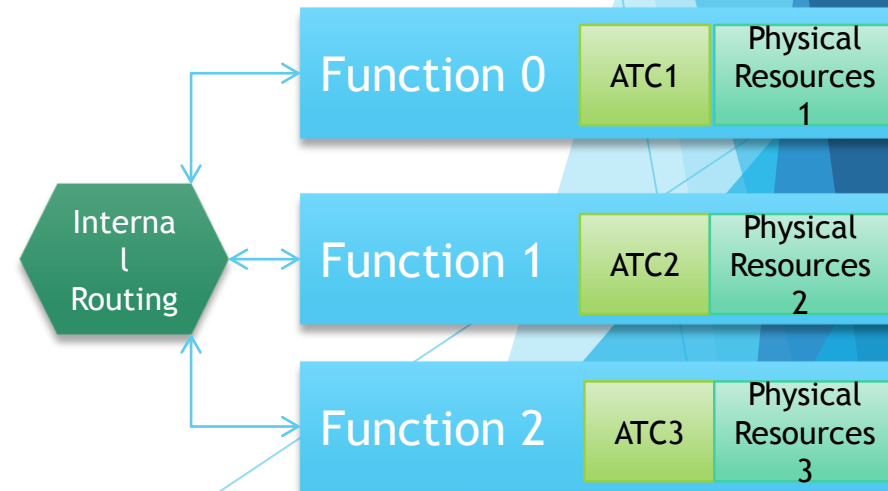
Components in PCIe Device

▶ Physical Resources

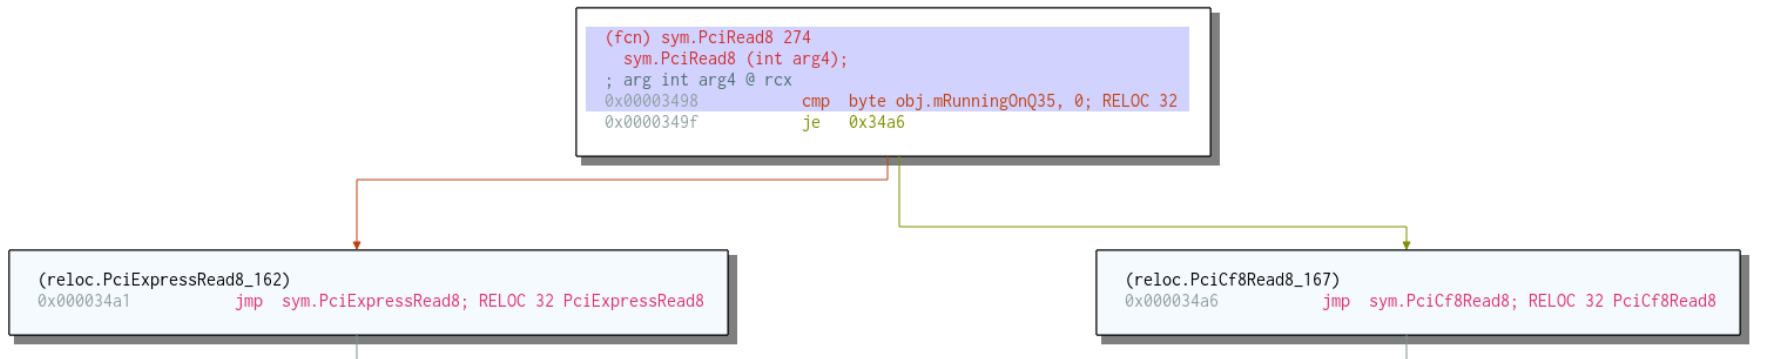
- ▶ Memory which allocated from physical memory

▶ ATC - Address Translation Cache

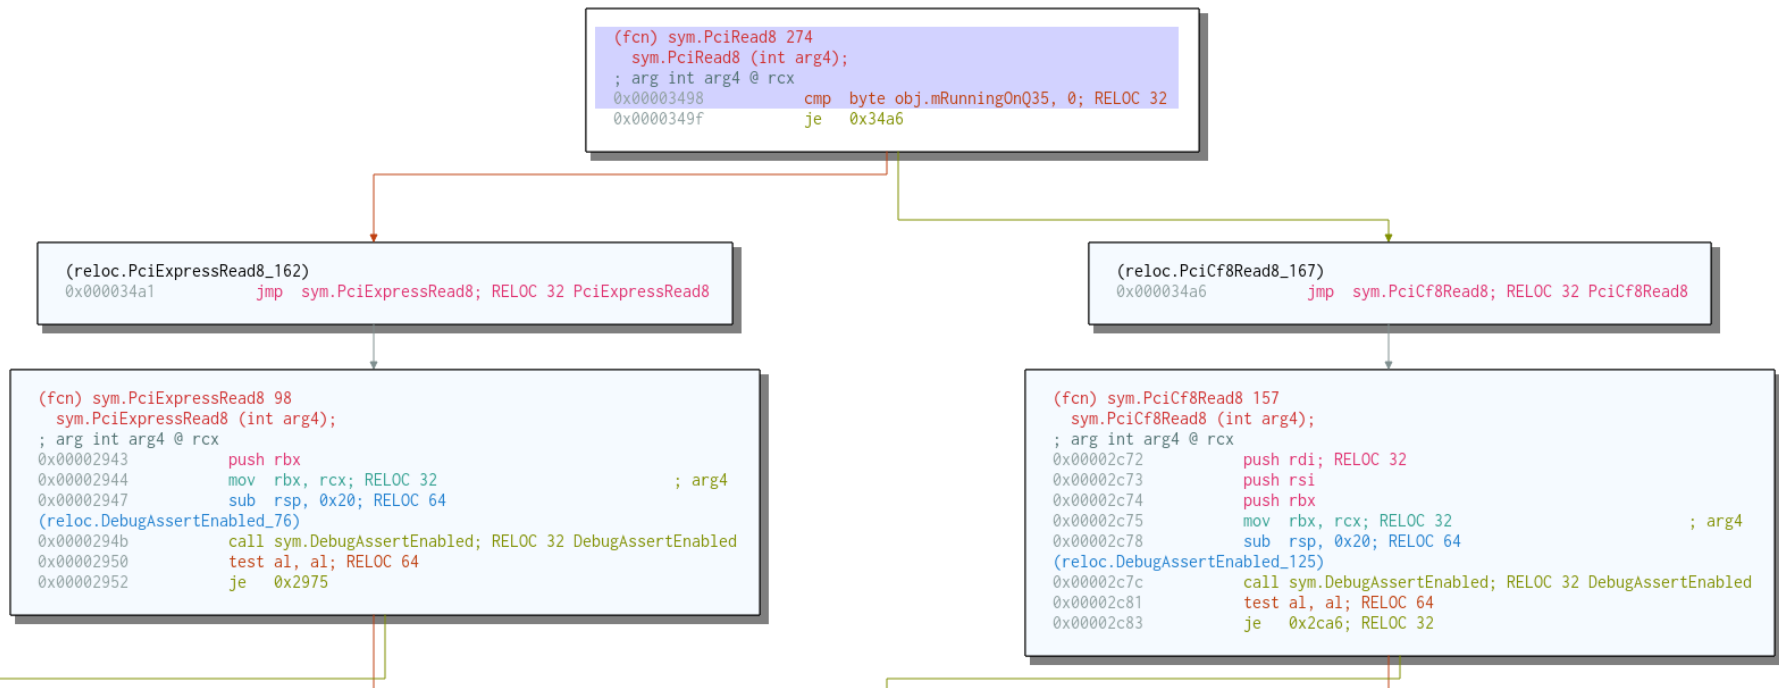
- ▶ A hardware stores recently used address translations.
- ▶ This term is used instead of TLB buffer
- ▶ To differentiate the **TLB used for I/O** from the **TLB used by the CPU**



Reversing PCI(e) device drivers - READ request



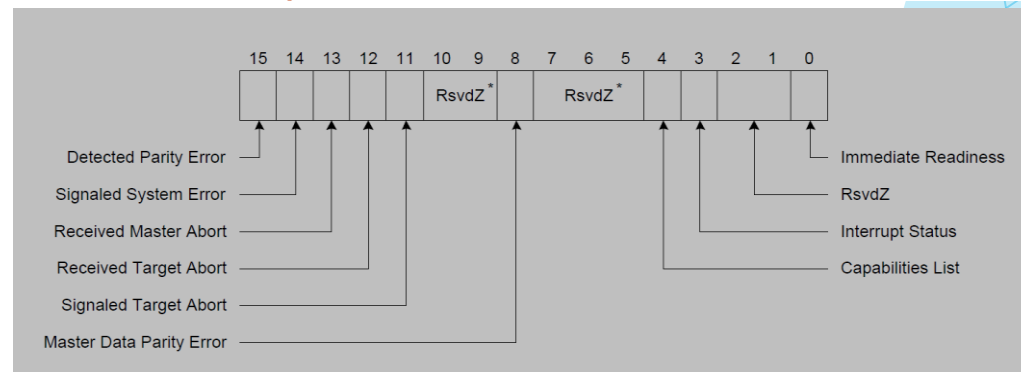
Reading PCIe devices vs Reading PCI devices



PCI Error? Set status register

```
(fcn) sym.PciCf8Read8 157
    sym.PciCf8Read8 (int arg4);
; arg int arg4 @ rcx
0x00002c72    push rdi; RELOC 32
0x00002c73    push rsi
0x00002c74    push rbx
0x00002c75    mov  rbx, rcx; RELOC 32                ; arg4
0x00002c78    sub  rsp, 0x20; RELOC 64
(reloc.DebugAssertEnabled_125)
0x00002c7c    call sym.DebugAssertEnabled; RELOC 32 DebugAssertEnabled
0x00002c81    test al, al; RELOC 64
0x00002c83    je   0x2ca6; RELOC 32
```

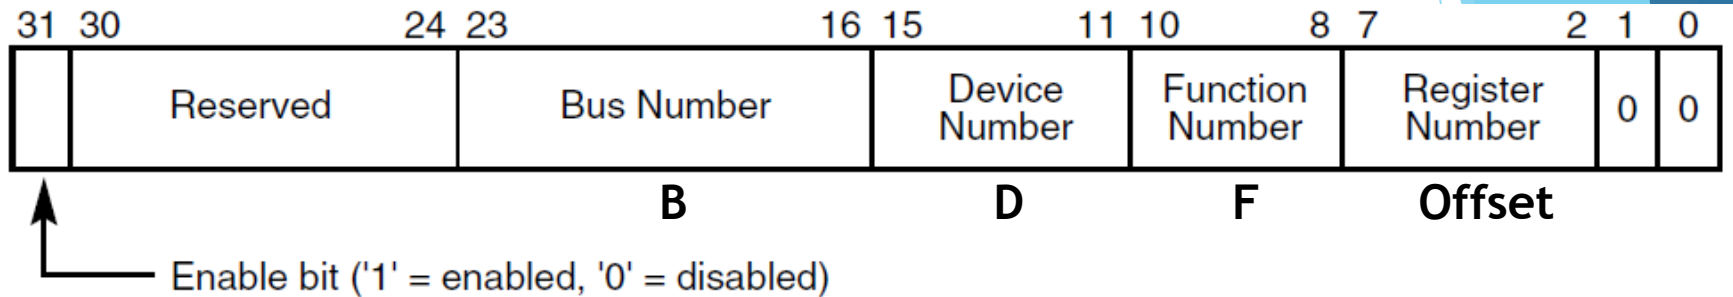
```
0x00002c85    test rbx, 0xffffffff0000f00; RELOC 32
0x00002c8c    je   0x2ca6; RELOC 32
```



Compatible PCI Configuration Space

- ▶ This refers to the software generation of PCI configuration transactions
 - ▶ those generated by the CPU/BIOS
- ▶ Compatible PCI provides 256 bytes of Configuration address space to the CPU/BIOS
- ▶ CPU/BIOS programs the registers contained therein to configure the device and system parameters
- ▶ Compatible PCI is configured using the port I/O address/data pair (CONFIG_ADDRESS, CONFIG_DATA)
- ▶ Two 32-bit I/O locations are used to generate configuration transactions
 - ▶ CF8h (CONFIG_ADDRESS)
 - ▶ CFCh (CONFIG_DATA)

I/O Port CONFIG_ADDRESS (CF8h)



- ▶ 32 bits (**GIMME BUS 0, DEVICE 31 (0x1F), Function 0, offset 0x88**)
- ▶ Port CF8h
- ▶ Bit 31 when set, all reads and writes to CONFIG_DATA are PCI Configuration transactions
- ▶ Bits 30:24 are read-only and must return 0 when read
- ▶ Bits 23:16 select a specific Bus in the system (up to 256 buses)
- ▶ Bits 15:11 specify a Device on the given Bus (up to 32 devices)
- ▶ Bits 10:8 Specify the function of a device (up to 8 devices)
- ▶ Bits 7:0 Select an offset within the Configuration Space (256 bytes max, DWORD-aligned as bits 1:0 are hard-coded 0)
- ▶ Addresses are often given in B/D/F, Offset notation (also written as B:D:F, Offset)

I/O Port CONFIG_DATA (CFCh)

- ▶ CONFIG_DATA can be accessed in DWORD, WORD, or BYTE configurations
- ▶ Reads and Writes to CONFIG_DATA with Bit 31 in CONFIG_ADDRESS set/enabled results in a PCI Configuration transaction to the device specified in CONFIG_ADDRESS
- ▶ PCI spec says that if Bit 31 is not enabled, then the transaction is forwarded out as Port I/O

Performing PCI device READ request

```
(reloc.SaveAndDisableInterrupts_167)
0x00002ca6      call sym.SaveAndDisableInterrupts; RELOC 32 SaveAndDisableInterrupts
0x00002cab      mov ecx, 0xcf8; RELOC 32
0x00002cb0      mov esi, eax; RELOC 32
(reloc.IoRead32_179)
0x00002cb2      call sym.IoRead32; RELOC 32 IoRead32
0x00002cb7      mov rdx, rbx
0x00002cba      mov ecx, 0xcf8
0x00002cbf      mov edi, eax
0x00002cc1      shr rdx, 4; RELOC 32
0x00002cc5      mov eax, ebx
0x00002cc7      and ebx, 3; RELOC 64
0x00002cca      and eax, 0xfc; RELOC 64
0x00002ccf      and edx, 0xffff00; RELOC 64
0x00002cd5      or  edx, eax
0x00002cd7      or  edx, 0x80000000; RELOC 64
(reloc.IoWrite32_222)
0x00002cdd      call sym.IoWrite32; RELOC 32 IoWrite32
0x00002ce2      lea rcx, [rbx + 0xcfc]; RELOC 64
(reloc.IoRead8_234)
0x00002ce9      call sym.IoRead8; RELOC 32 IoRead8
0x00002cee      mov edx, edi
0x00002cf0      mov ecx, 0xcf8; RELOC 64
0x00002cf5      mov ebx, eax
(reloc.IoWrite32_248)
0x00002cf7      call sym.IoWrite32; RELOC 32 IoWrite32
0x00002cfc      movzx ecx, sil
(reloc.SetInterruptState_1)
0x00002d00      call sym.SetInterruptState; RELOC 32 SetInterruptState
0x00002d05      add rsp, 0x20
0x00002d09      mov eax, ebx
0x00002d0b      pop rbx
0x00002d0c      pop rsi; RELOC 32
0x00002d0d      pop rdi
0x00002d0e      ret
```

PCI Express - READ example

```
(fcn) sym.PciExpressRead8 98
  sym.PciExpressRead8 (int arg4);
; arg int arg4 @ rcx
0x00002943      push rbx
0x00002944      mov  rbx, rcx; RELOC 32          ; arg4
0x00002947      sub  rsp, 0x20; RELOC 64
(reloc.DebugAssertEnabled_76)
0x0000294b      call sym.DebugAssertEnabled; RELOC 32 DebugAssertEnabled
0x00002950      test al, al; RELOC 64
0x00002952      je   0x2975
```

The diagram shows a flow from a function signature to a code block. A grey arrow points from the top to the first code block. A red arrow points from the bottom of the first code block to the second code block. A green arrow points from the bottom of the second code block to the right.


```
0x00002954      test rbx, 0xffffffff00000000; RELOC 64
0x0000295b      je   0x2975; RELOC 32
```


If everything is OK, perform MmioRead8()



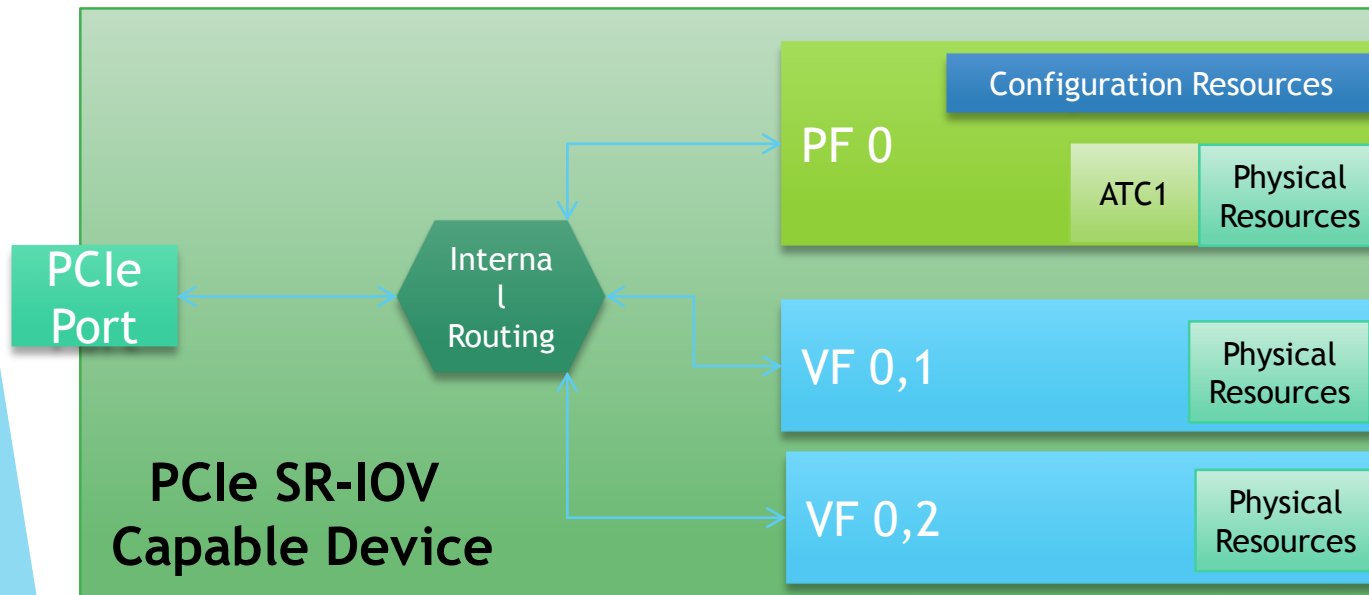
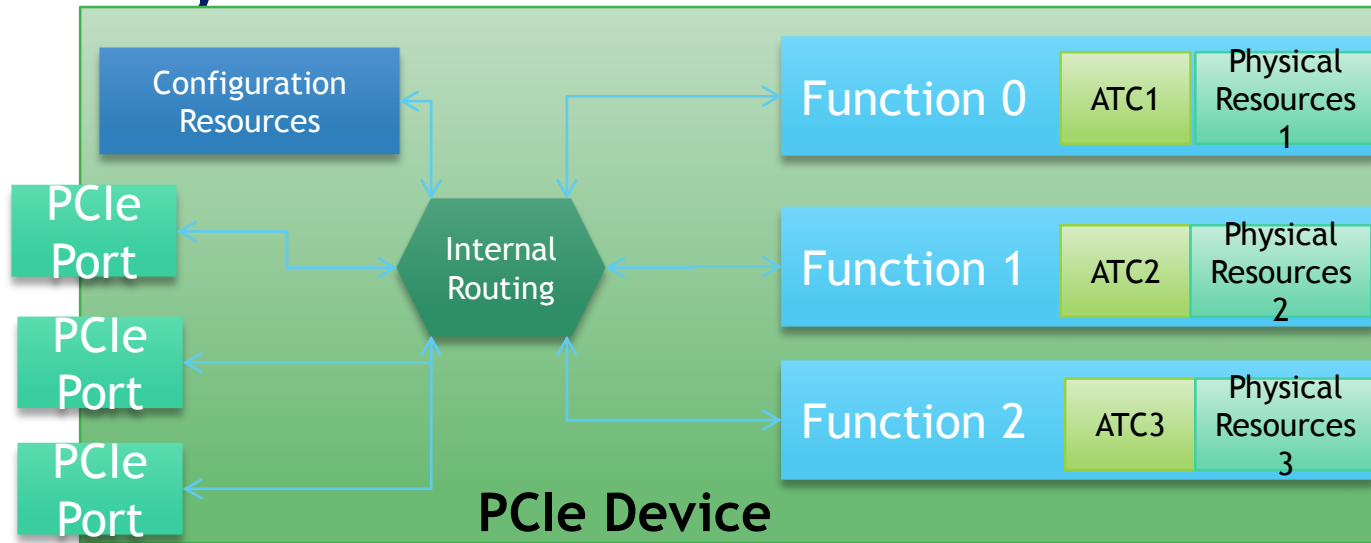
```
(reloc._gPcd_FixedAtBuild_PcdPciExpressBaseAddress_120)
0x00002975      add  rbx, qword [obj._gPcd_FixedAtBuild_PcdPciExpressBaseAddress]; RELOC 32 ...
0x0000297c      add  rsp, 0x20; RELOC 32
0x00002980      mov  rcx, rbx; RELOC 64
0x00002983      pop  rbx; RELOC 32
(reloc.MmioRead8_133)
0x00002984      jmp  sym.MmioRead8; RELOC 32 MmioRead8
```

RCX is transformed as a Base+Offset



```
(fcn) sym.MmioRead8 28
    sym.MmioRead8 (int arg4);
; arg int arg4 @ rcx
0x00002106    push rbx
0x00002107    mov  rbx, rcx                                ; arg4
0x0000210a    sub  rsp, 0x20; RELOC 32
(reloc.MemoryFence_15)
0x0000210e    call sym.MemoryFence; RELOC 32 MemoryFence
0x00002113    mov  bl, byte [rbx]; RELOC 64
(reloc.MemoryFence_22)
0x00002115    call sym.MemoryFence; RELOC 32 MemoryFence
0x0000211a    add  rsp, 0x20; RELOC 32
0x0000211e    mov  eax, ebx
0x00002120    pop  rbx; RELOC 64
0x00002121    ret
```

Physical V.S. Virtual



PCIe SR-IOV Capable Device

SR-IOV

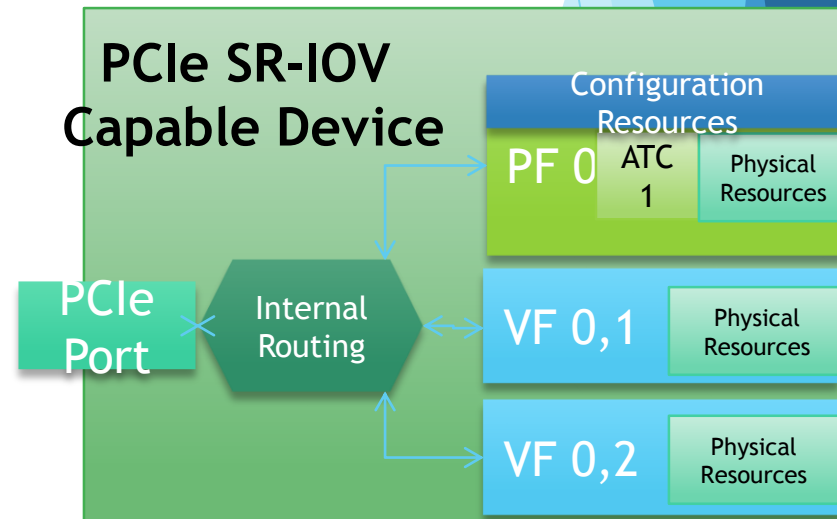
- ▶ A technique performs and manages PCIe Virtualization.

PF - physical Function

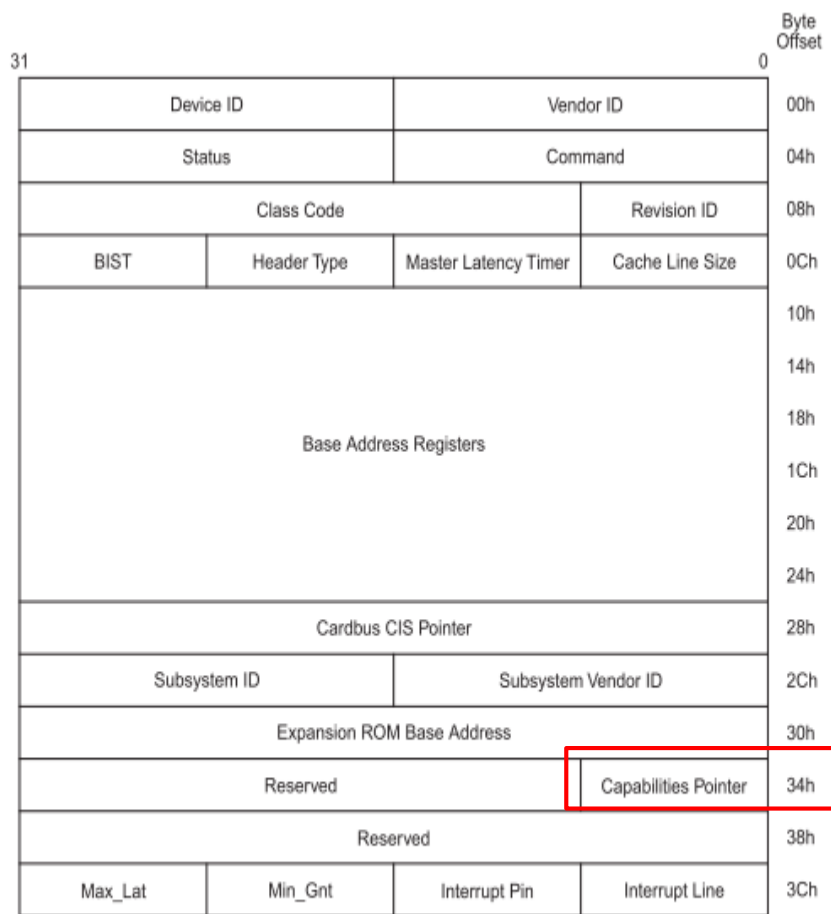
- ▶ Provide full PCIe functionality, including the SR-IOV capabilities
- ▶ Discover the **page sizes** supported by a PF and its associated VF

VF - virtual Function

- ▶ A **“light-weight”** PCIe function that is **directly accessible by an SI**, including an isolated memory space, a work queue, interrupts and command processing.
- ▶ **For data movement**
- ▶ Can be optionally migrated from one PF to another PF
- ▶ Can be serially shared by different SI




Extended Capabilities



OM14316

Standard PCI Configuration
Initial Header Type 0 with
Capability pointer at Offset 34h



A-0318

Figure 3-2: Type 0 Configuration Space Header

Figure 3-2: Capabilities Linked List

PCI Express Capability Structure

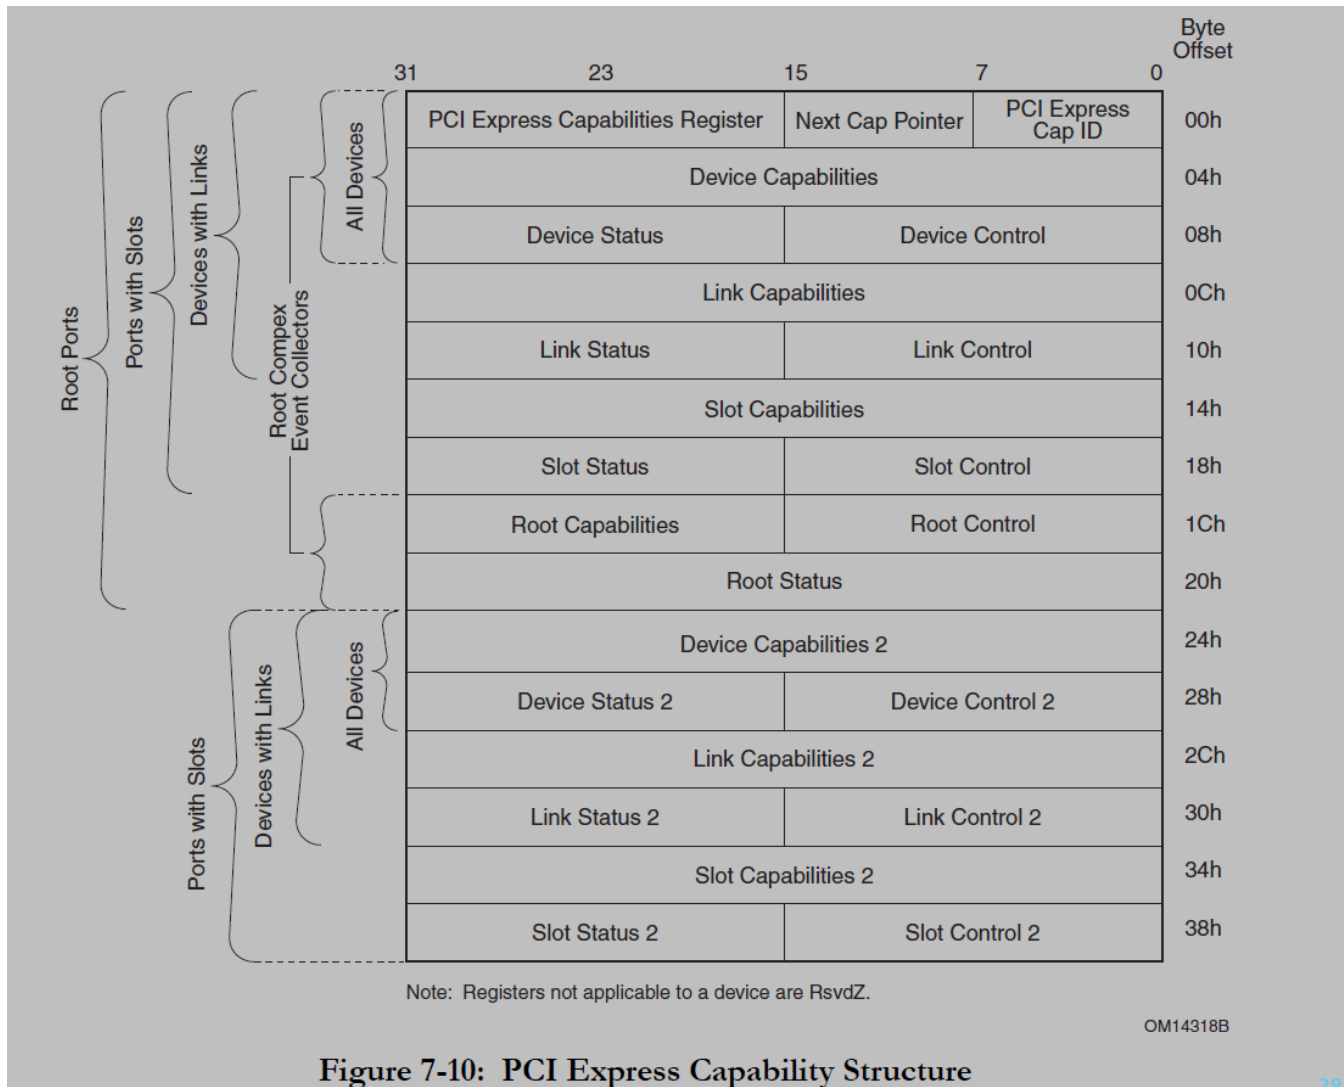
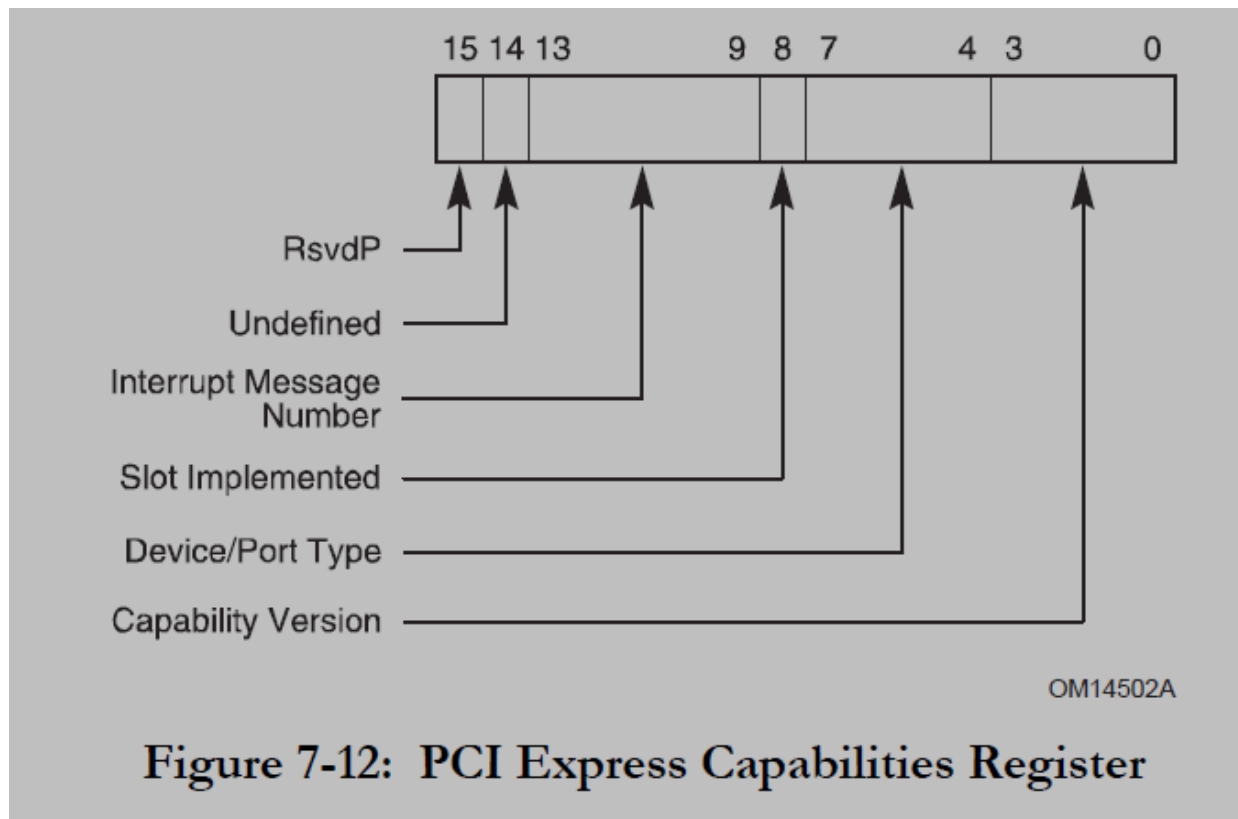


Figure 7-10: PCI Express Capability Structure

PCI Express Capabilities Register



PCI Express Device/Port TYPE

Bit Location	Register Description	Attributes																		
7:4	<p>Device/Port Type – Indicates the specific type of this PCI Express Function. Note that different Functions in a multi-Function device can generally be of different types.</p> <p>Defined encodings are:</p> <table><tr><td>0000b</td><td>PCI Express Endpoint</td></tr><tr><td>0001b</td><td>Legacy PCI Express Endpoint</td></tr><tr><td>0100b</td><td>Root Port of PCI Express Root Complex*</td></tr><tr><td>0101b</td><td>Upstream Port of PCI Express Switch*</td></tr><tr><td>0110b</td><td>Downstream Port of PCI Express Switch*</td></tr><tr><td>0111b</td><td>PCI Express to PCI/PCI-X Bridge*</td></tr><tr><td>1000b</td><td>PCI/PCI-X to PCI Express Bridge*</td></tr><tr><td>1001b</td><td>Root Complex Integrated Endpoint</td></tr><tr><td>1010b</td><td>Root Complex Event Collector</td></tr></table> <p>*This value is only valid for Functions that implement a Type 01h PCI Configuration Space header.</p> <p>All other encodings are Reserved.</p> <p>Note that the different Endpoint types have notably different requirements in Section 1.3.2 regarding I/O resources, Extended Configuration Space, and other capabilities.</p>	0000b	PCI Express Endpoint	0001b	Legacy PCI Express Endpoint	0100b	Root Port of PCI Express Root Complex*	0101b	Upstream Port of PCI Express Switch*	0110b	Downstream Port of PCI Express Switch*	0111b	PCI Express to PCI/PCI-X Bridge*	1000b	PCI/PCI-X to PCI Express Bridge*	1001b	Root Complex Integrated Endpoint	1010b	Root Complex Event Collector	RO
0000b	PCI Express Endpoint																			
0001b	Legacy PCI Express Endpoint																			
0100b	Root Port of PCI Express Root Complex*																			
0101b	Upstream Port of PCI Express Switch*																			
0110b	Downstream Port of PCI Express Switch*																			
0111b	PCI Express to PCI/PCI-X Bridge*																			
1000b	PCI/PCI-X to PCI Express Bridge*																			
1001b	Root Complex Integrated Endpoint																			
1010b	Root Complex Event Collector																			

SR-IOV Extended Capabilities

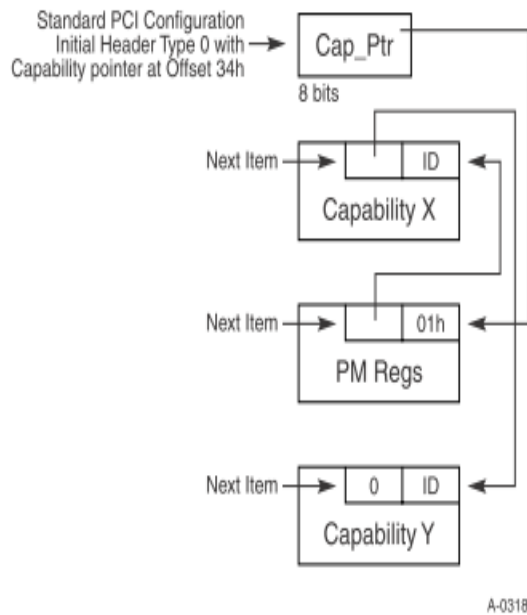
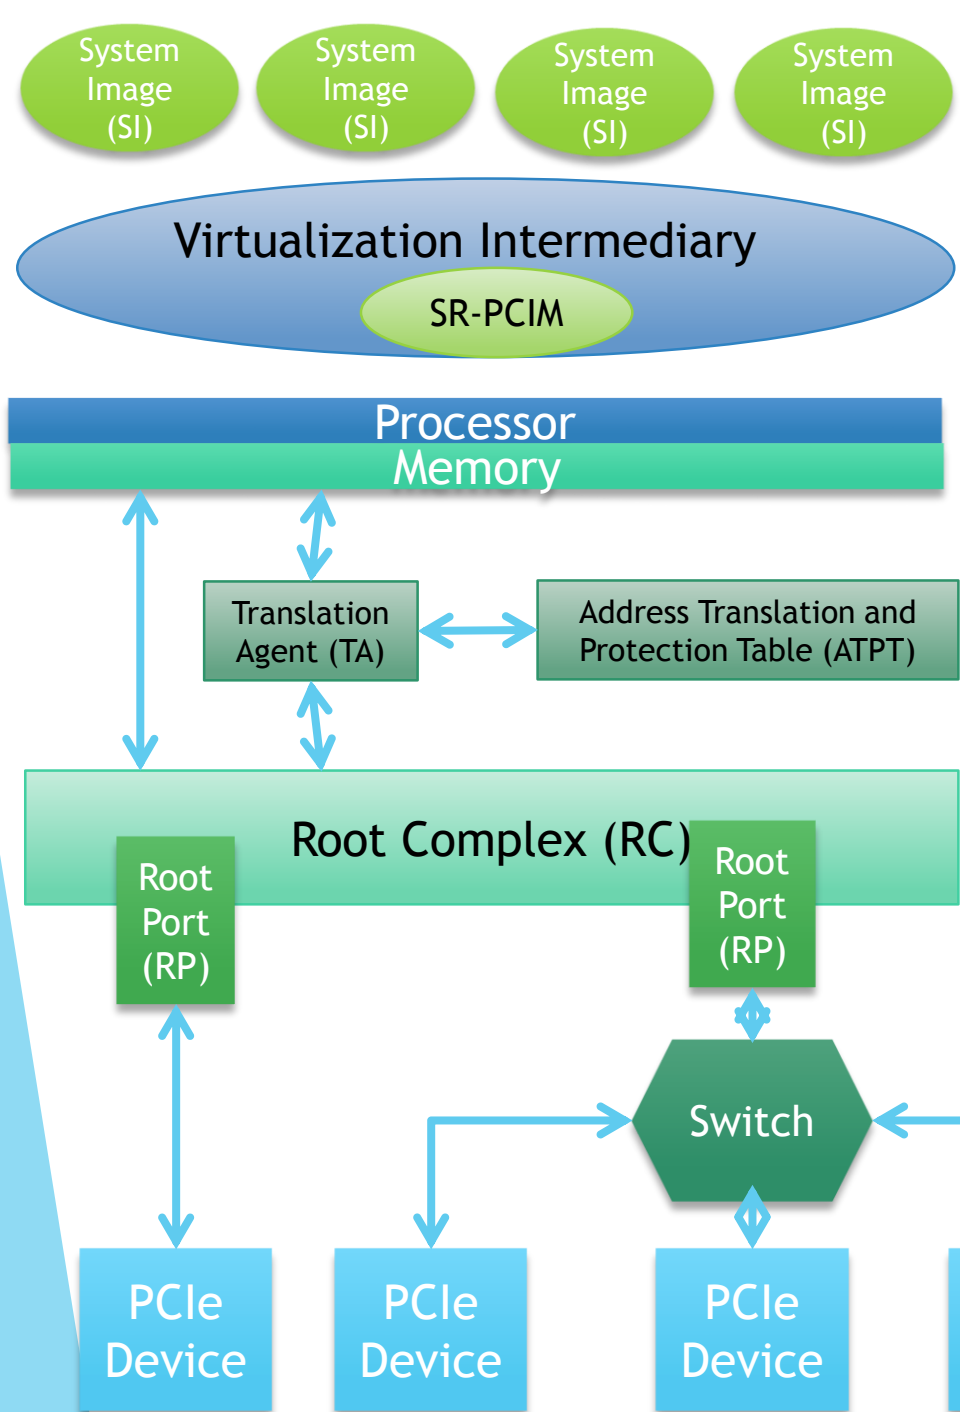


Figure 3-2: Capabilities Linked List

31	24	23	20	19	16	15	0	Byte Offset
Next Capability Offset			Capability Version		PCI Express Extended Capability ID			00h
SR-IOV Capabilities								04h
SR-IOV Status				SR-IOV Control				08h
TotalVFs (RO)				InitialVFs (RO)				0Ch
RsvdP		Function Dependency Link (RO)		NumVFs (RW)				10h
VF Stride (RO)				First VF Offset (RO)				14h
VF Device ID (RO)				RsvdP				18h
Supported Page Sizes (RO)								1Ch
System Page Size (RW)								20h
VF BAR0 (RW)								24h
VF BAR1 (RW)								28h
VF BAR2 (RW)								2Ch
VF BAR3 (RW)								30h
VF BAR4 (RW)								34h
VF BAR5 (RW)								38h
VF Migration State Array Offset (RO)								3Ch

A-0634A

Figure 3-1: Single Root I/O Virtualization Extended Capabilities



▶ SR-PCIM

- ▶ Configure SR-IOV Capability
- ▶ Management of PFs and VFs
- ▶ Processing of error events
- ▶ Device controls
 - ▶ Power management
 - ▶ Hot-plug

SR-PCIM

Components of SR-IOV

- ▶ TA - Translation Agent
 - ▶ Translate address within a PCIe transaction into the associated platform physical address.
 - ▶ Hardware or combination of hardware and software
 - ▶ A TA may also support to enable a PCIe function to obtain address translations a priori to DMA access to the associated memory.



Components of SR-IOV

- ▶ ATPT - Address Translation and Protection Table
 - ▶ Contain the set of address translations accessed by a TA to Process PCIe requests
 - ▶ DMA Read/Write
 - ▶ Interrupt requests
 - ▶ DMA Read/Write requests are translated through a combination of the **Routing ID** and **the address contained within a PCIe transaction**
 - ▶ In PCIe, interrupts are treated as memory write operations.
 - ▶ Though the combination of the Routing ID and the address contained within a PCIe transaction as well



ARI - Alternative Routing ID Interpretation

- ▶ Routing ID is used to forward requests to the corresponding PFs and VFs
- ▶ All VFs and PFs must have distinct Routing IDs
- ▶ ARI provides a mechanism to allow single PCIe component to support up to 256 functions.
 - ▶ Originally there are 8 functions at most in a PCIe.

15 ----- 8	7-----3	2-----0
Bus #	Device #	Function #

Table 1. Traditional Routing

15 ----- 8	7 ----- 0
Bus #	Identifier

Table 2. Alternate Routing

Figure from Intel PCI-SIG SR_IOV p

ARI - Alternative Routing ID Interpretation

Table 2-1: VF Routing ID Algorithm

VF Number	VF Routing ID
VF 1	$(\text{PF Routing ID} + \text{First VF Offset}) \text{ Modulo } 2^{16}$
VF 2	$(\text{PF Routing ID} + \text{First VF Offset} + \text{VF Stride}) \text{ Modulo } 2^{16}$
VF 3	$(\text{PF Routing ID} + \text{First VF Offset} + 2 * \text{VF Stride}) \text{ Modulo } 2^{16}$
...	...
VF N	$(\text{PF Routing ID} + \text{First VF Offset} + (N-1) * \text{VF Stride}) \text{ Modulo } 2^{16}$
...	...
VF NumVFs (last one)	$(\text{PF Routing ID} + \text{First VF Offset} + (\text{NumVFs}-1) * \text{VF Stride}) \text{ Modulo } 2^{16}$

Figure from SR-IOV Specification revision ...




Figure from Intel PCI-SIG SR_IOV prim

ATS - Address Translation Services

- ▶ ATS provides a mechanism allowing a virtual machine to perform **DMA** transaction directly to and from a PCIe endpoint.

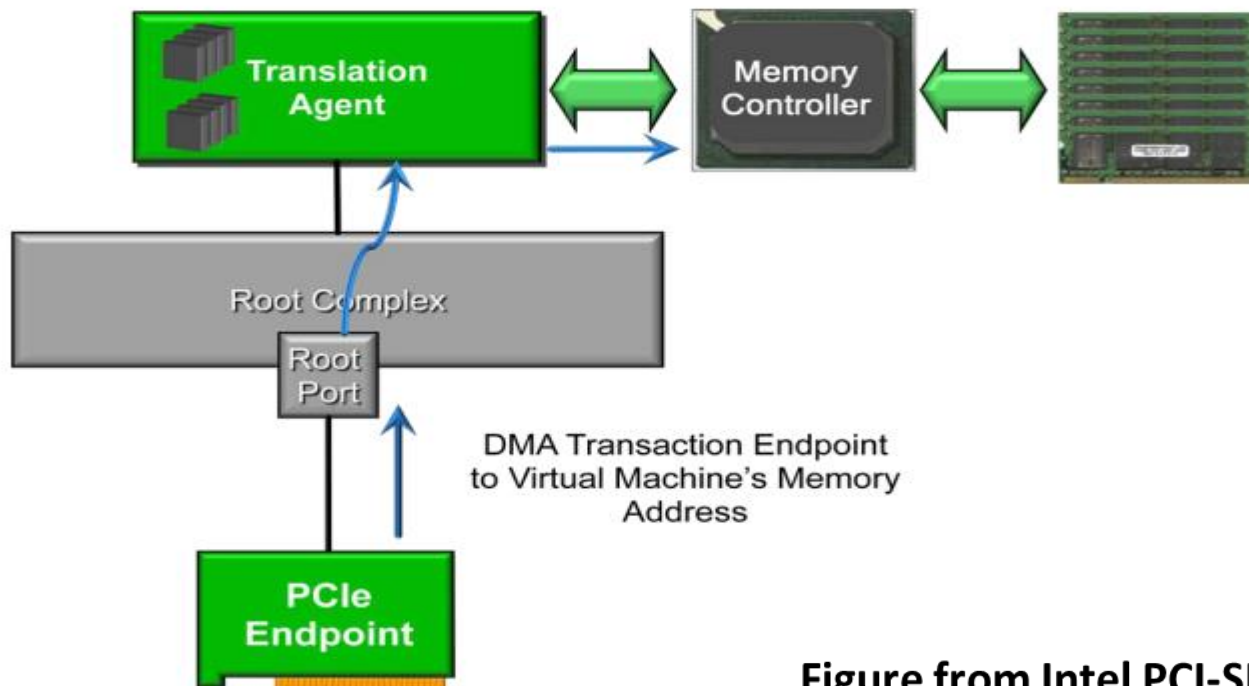
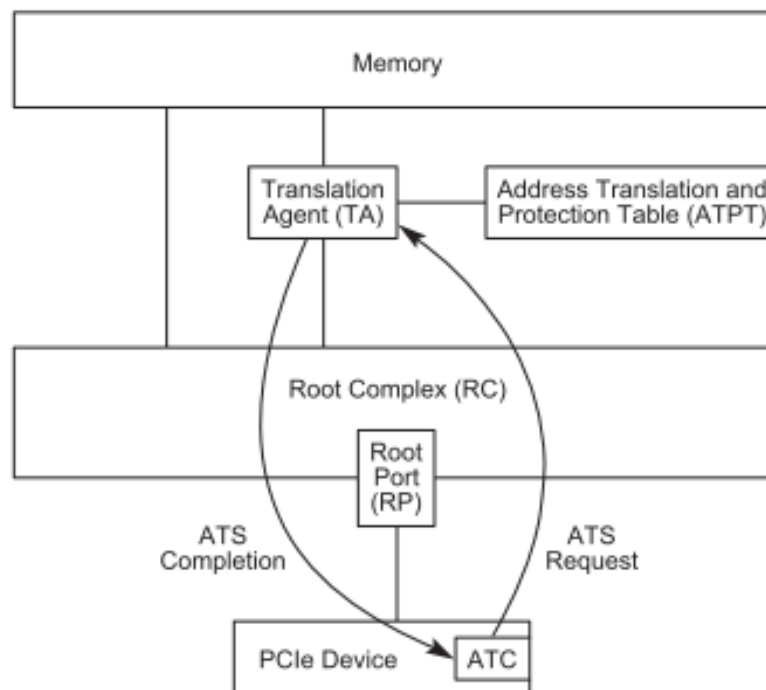


Figure from Intel PCI-SIG SR_IOV prim

ATS - Address Translation Services

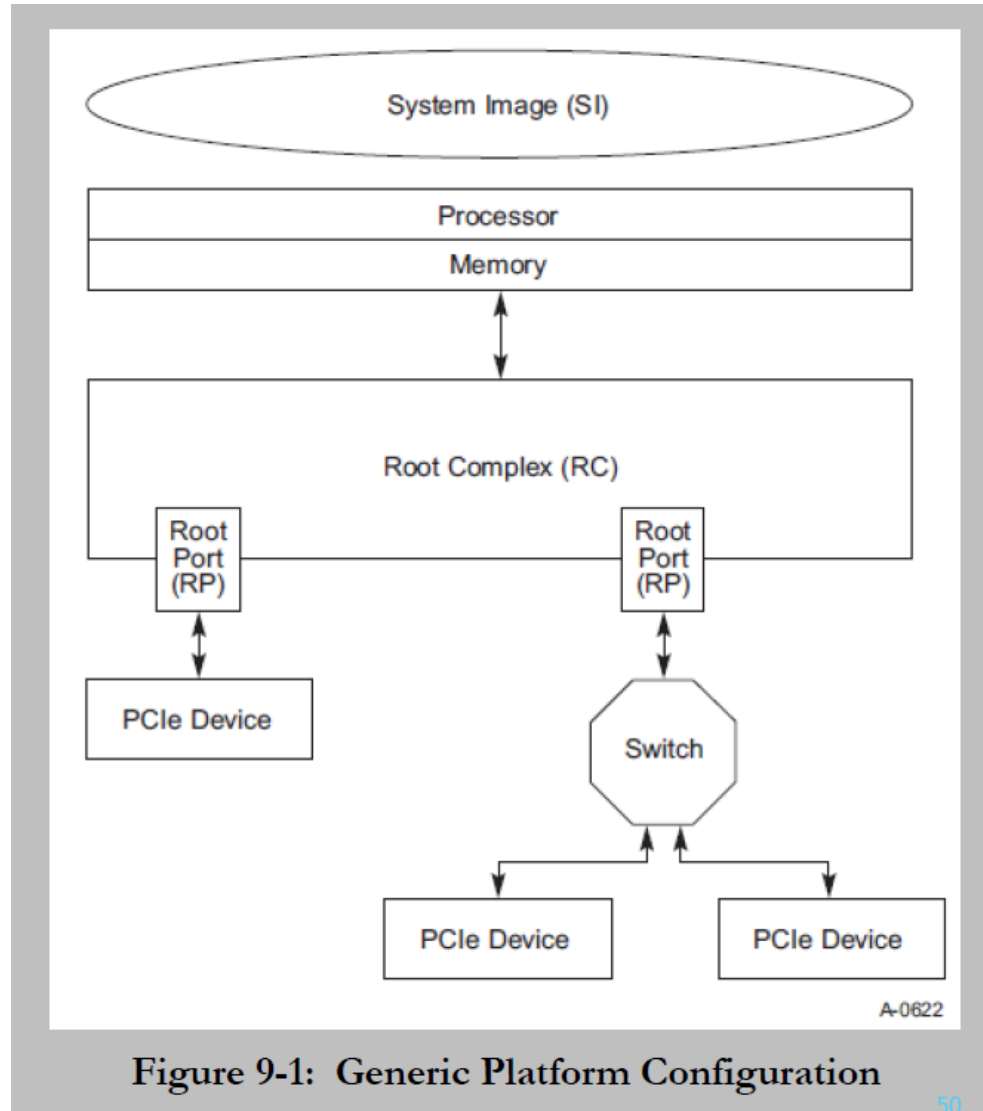
- ▶ ATS uses a **request-completion** protocol between a Device and a Root Complex (RC)



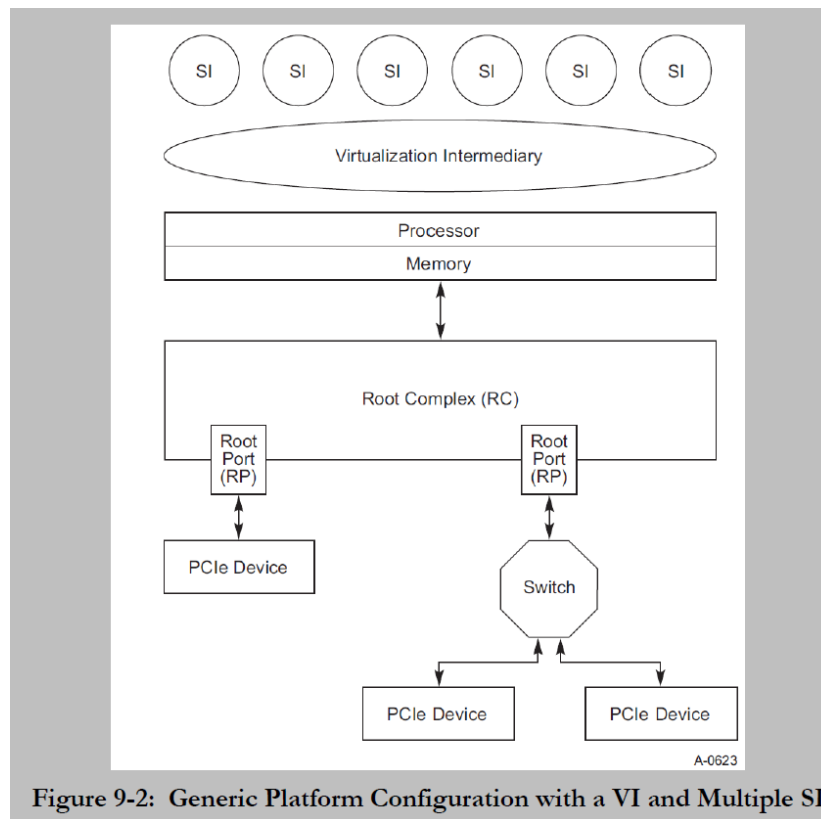
ATS - Address Translation Services

- ▶ When the Function receives the ATS Translation Completion
 - ▶ Either updates its ATC to reflect the translation
 - ▶ Or notes that a translation does not exist.
- ▶ The Function generates subsequent requests using
 - ▶ Either a translated address
 - ▶ Or an un-translated address based on the results of the Completion.

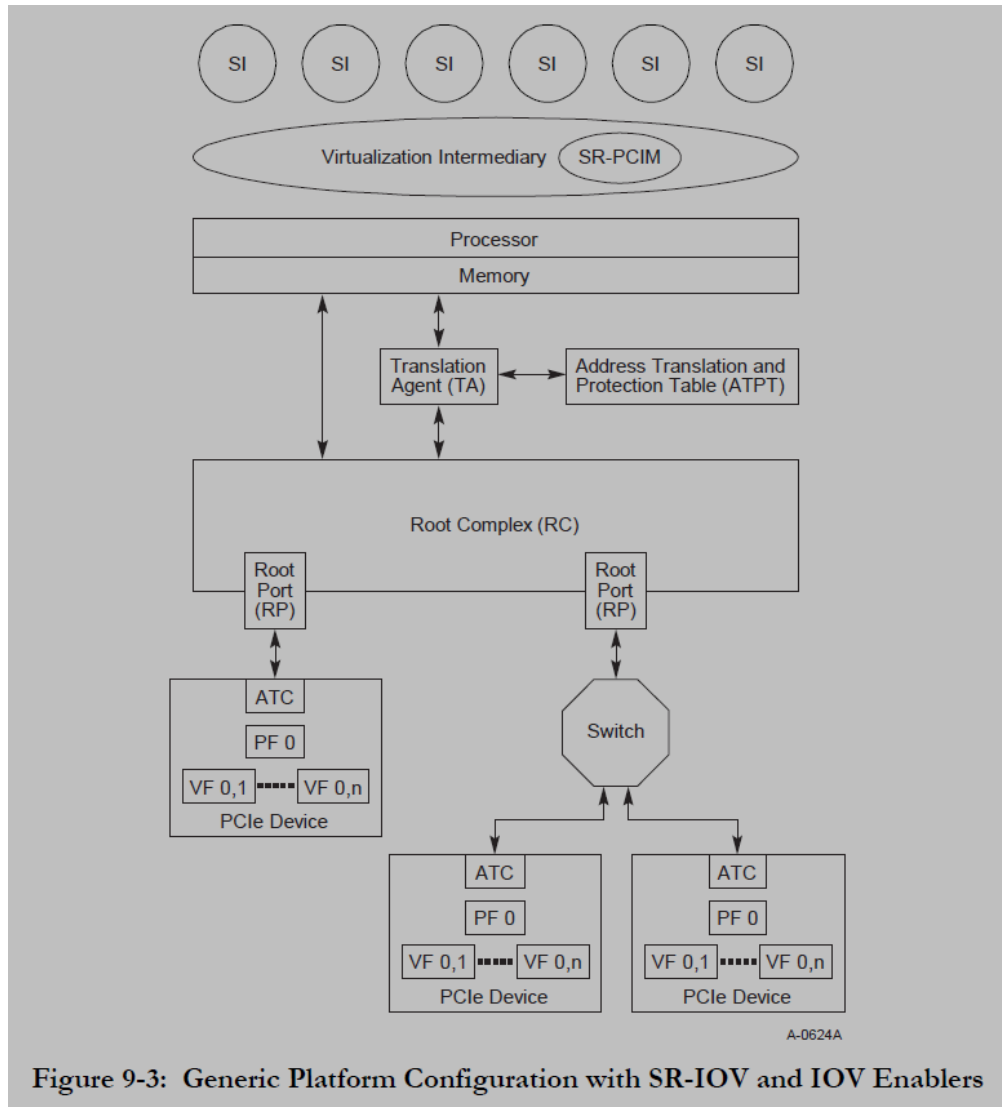
“Vanilla” PCIe Device



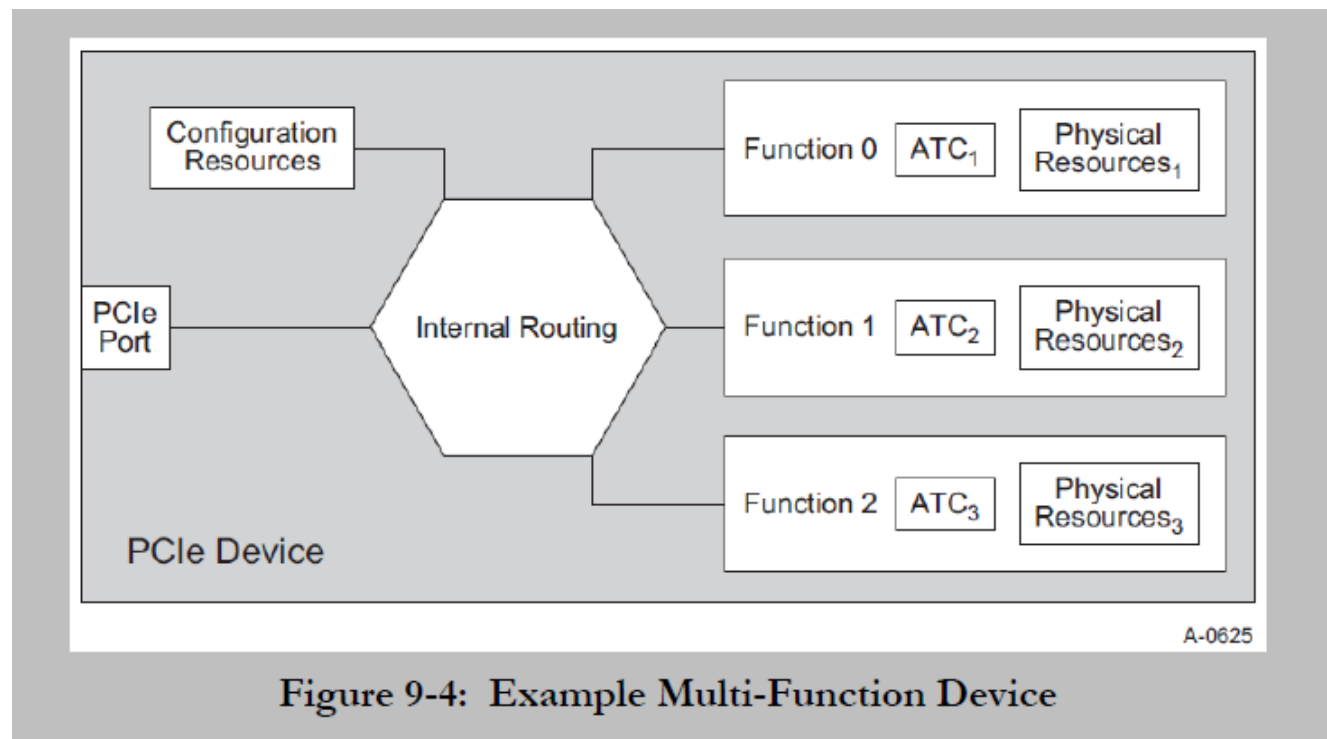
PCI Express Fabric with PCIe Express Physical Functions



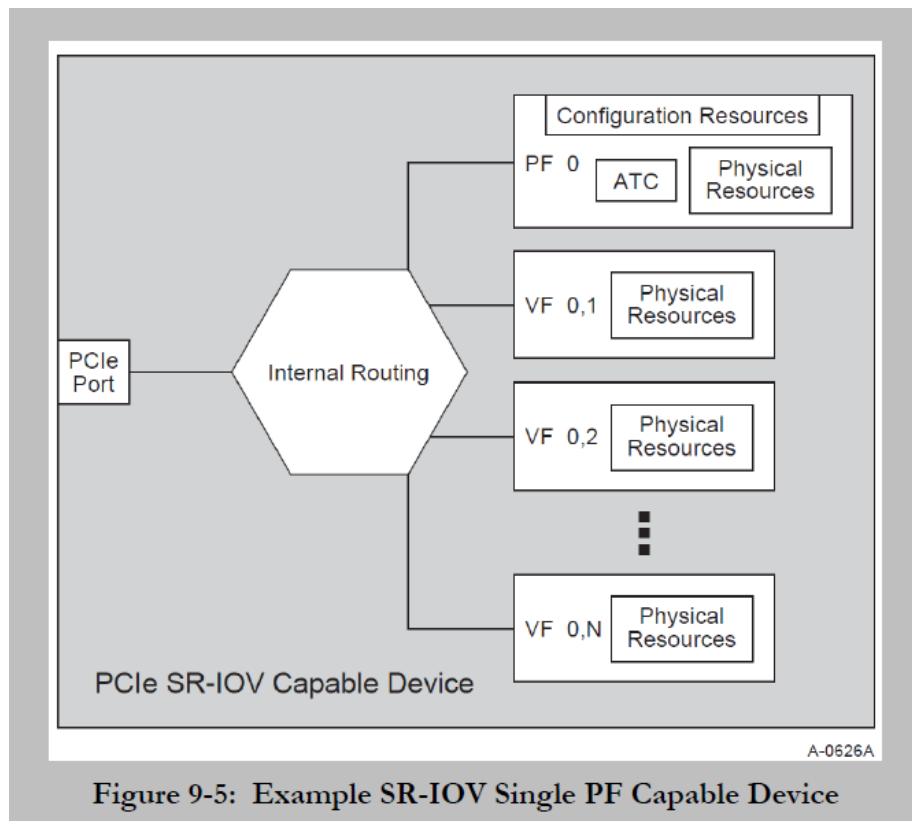
PCI Express with SR-IOV enabled



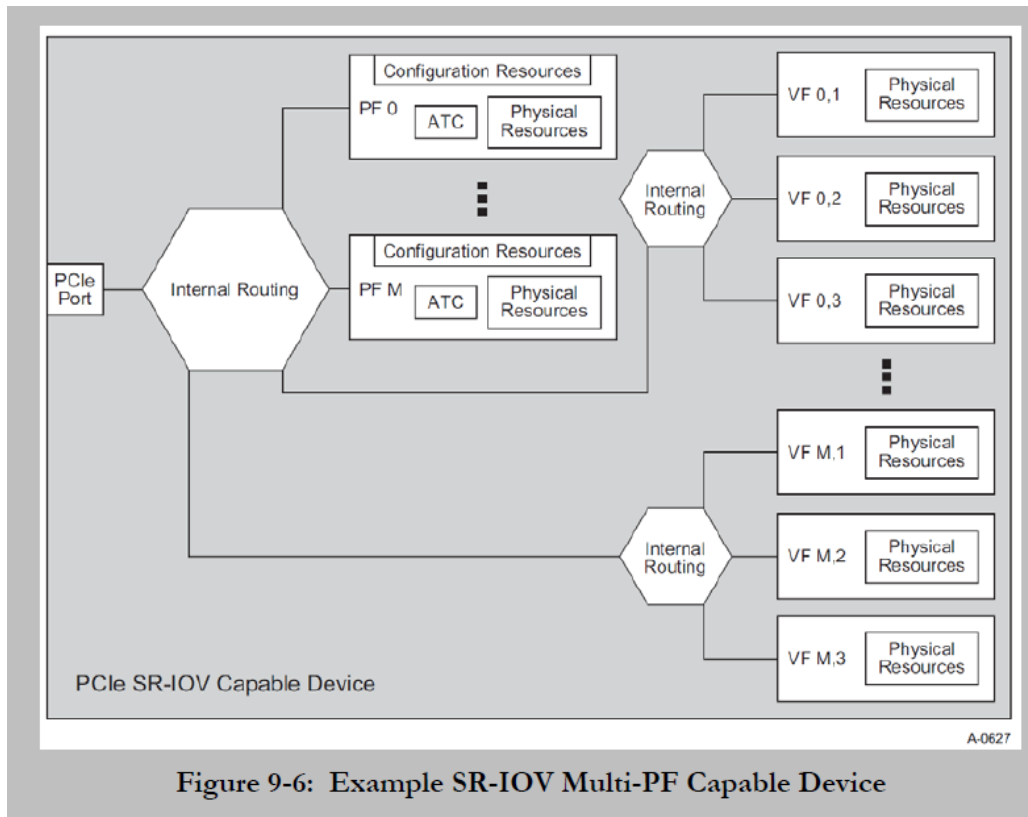
PCI Express device with Multiple physical functions



PCI Express device, SR-IOV capable with single PF (our focus)



PCI Express device, SR-IOV capable with Multiple PFs

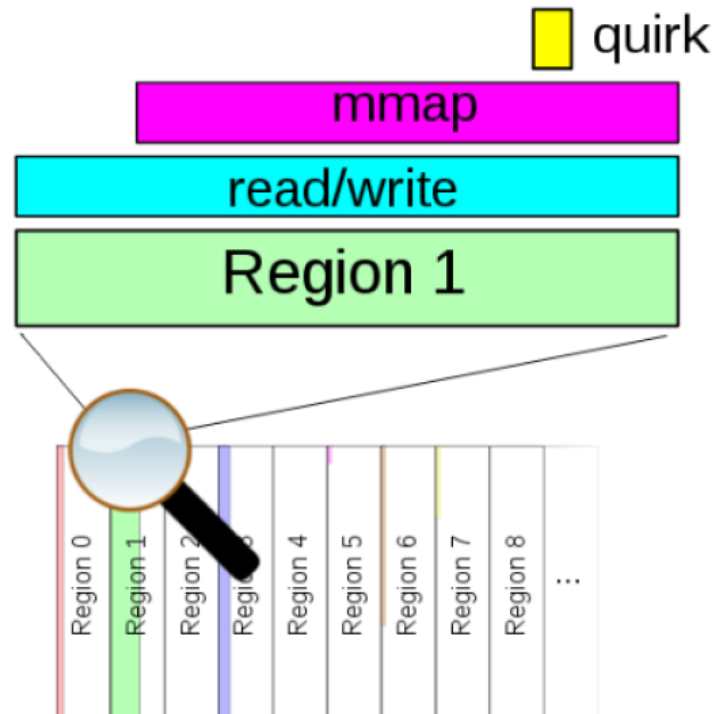


SR-IOV Configuration Space Mapping

- ❑ Determine desired number of Virtual Functions from *InitialVFs* field
- ❑ Program *NumVFs* field to match
- ❑ Multi-Root adds a further layer where configuration software **first** allocates VFs to Virtual Hierarchies – thus *InitialVFs* may be less than *TotalVFs*

31 24 23 20 19 16 15			0	Byte Offset
Next Capability Offset			Capability Version	PCI Express Extended Capability ID
			00h	
			SR-IOV Capabilities	
			04h	
			SR-IOV Status	
			SR-IOV Control	
			08h	
			TotalVFs (RO)	
			InitialVFs (RO)	
			0Ch	
			RsvdP	
			Function Dependency Link (RO)	
			NumVFs (RW)	
			10h	
			VF Stride (RO)	
			First VF Offset (RO)	
			14h	
			VF Device ID (RO)	
			RsvdP	
			18h	
			Supported Page Sizes (RO)	
			1Ch	
			System Page Size (RW)	
			20h	
			VF BAR0 (RW)	
			24h	
			VF BAR1 (RW)	
			28h	
			VF BAR2 (RW)	
			2Ch	
			VF BAR3 (RW)	
			30h	
			VF BAR4 (RW)	
			34h	
			VF BAR5 (RW)	
			38h	
			VF Migration State Array Offset (RO)	
			3Ch	

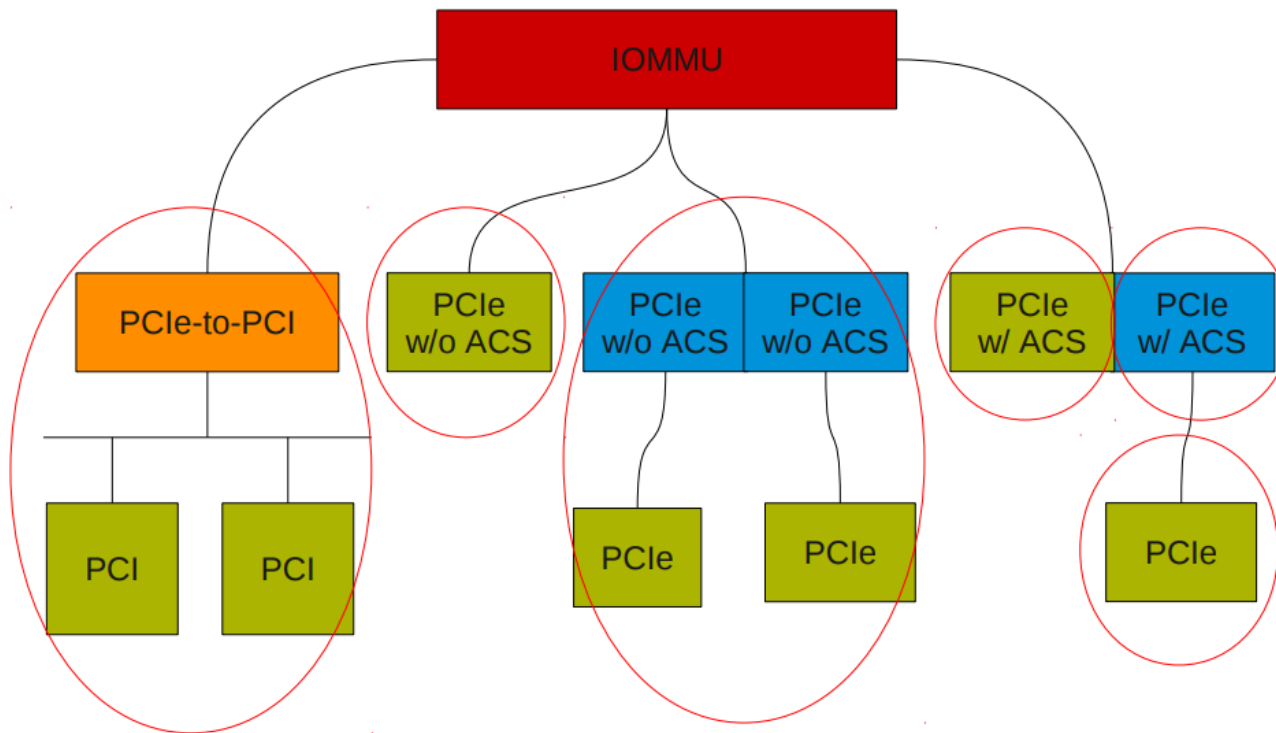
The device implements SR-IOV capabilities using vfio




VFIO - User-space Virtual Function I/O

- ▶ A new user level driver framework for Linux
- ▶ • Originally developed by Tom Lyon (Cisco)
- ▶ • IOMMU-based DMA and interrupt isolation
- ▶ • Full devices access (MMIO, I/O port, PCI config)
- ▶ • Efficient interrupt mechanisms
- ▶ • Modular IOMMU and device backends

IOMMU is in charge of “isolation GROUPS”



VFIO in a nutshell



Reverse engineering SR-IOV process - Probe VFs of a PF

```
(fcn) sym.vfio_pci_probe 333
    sym.vfio_pci_probe (int arg_58h);
; arg int arg_58h @ rbp+0x58
0x080013c0      call 0x80013c5; RELOC 32 __fentry__
0x080013c5      push r14
0x080013c7      push r13
0x080013c9      push r12
0x080013cb      push rbp
0x080013cc      push rbx
0x080013cd      cmp byte [rdi + 0x49], 0
0x080013d1      jne 0x80014ec
```

```
0x080013d7      lea r12, [rdi + 0xa0] ; 160
0x080013de      mov rbx, rdi
0x080013e1      mov rdi, r12
(reloc.vfio_iommu_group_get)
0x080013e4      call 0x80013e9; RELOC 32 vfio_iommu_group_get
0x080013e9      test rax, rax
0x080013ec      mov r14, rax
0x080013ef      je 0x80014ec
```

If the device includes SR-IOV capability, trace functions

```
0x080013d7    lea r12, [rdi + 0xa0]                ; 160
0x080013de    mov rbx, rdi
0x080013e1    mov rdi, r12
(reloc.vfio_iommu_group_get)
0x080013e4    call 0x80013e9; RELOC 32 vfio_iommu_group_get
0x080013e9    test rax, rax
0x080013ec    mov r14, rax
0x080013ef    je 0x80014ec
```

```
0x080013f5    mov rdi, qword [0x080013fc]                ; [0x80013fc:8]=0x80c0be000...
0x080013fc    mov edx, 0x100                          ; 256
0x08001401    mov esi, 0x14080c0
0x08001406    call 0x800140b; RELOC 32 kmem_cache_alloc_trace
0x0800140b    test rax, rax
0x0800140e    mov rbp, rax
0x08001411    je 0x80014f7
```

```
0x080014ec    mov r13d, 0xfffffffffa                  ; 4294967274
0x080014f2    jmp 0x8001471
```

```
0x08001471    pop rbx
0x08001472    mov eax, r13d
0x08001475    pop rbp
0x08001476    pop r12
0x08001478    pop r13
0x0800147a    pop r14
0x0800147c    ret
```

If we find a VF, add it to the group (according to IOMMU)



Check for special capabilities (VGA Passthrough, etc.)

↓

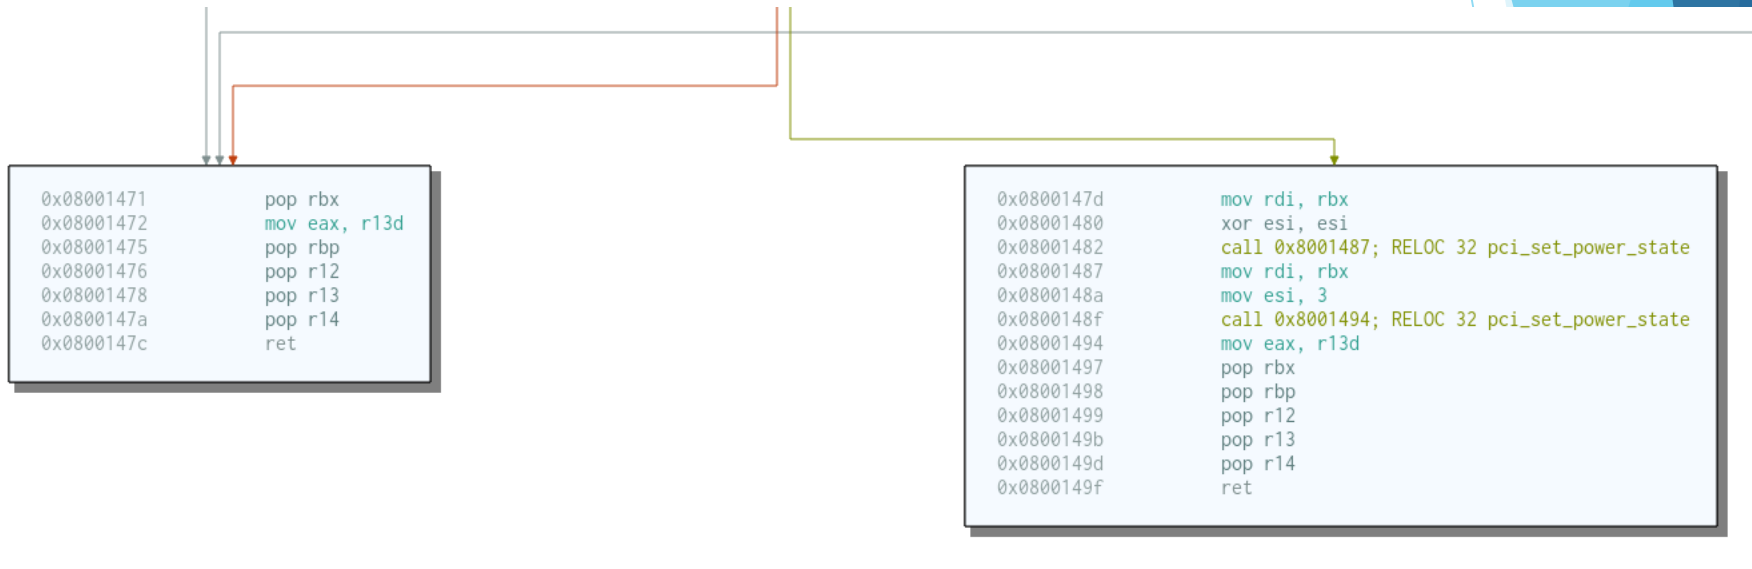
```
0x0800145b    mov eax, dword [rbx + 0x44]           ; [0x44:4]=-1 ; 'D' ; 68
0x0800145e    shr eax, 8
0x08001461    cmp eax, 0x300                      ; 768
0x08001466    je 0x80014bf
```

L

↓

```
0x080014bf    mov rcx, 0; RELOC 32
0x080014c6    xor edx, edx
0x080014c8    mov rsi, rbp
0x080014cb    mov rdi, rbx
(reloc.vga_client_register)
0x080014ce    call 0x80014d3; RELOC 32 vga_client_register
0x080014d3    xor esi, esi
0x080014d5    mov rdi, rbp
0x080014d8    call sym.vfio_pci_set_vga_decode
0x080014dd    mov rdi, rbx
0x080014e0    mov esi, eax
(reloc.vga_set_legacy_decoding)
0x080014e2    call 0x80014e7; RELOC 32 vga_set_legacy_decoding
0x080014e7    jmp 0x8001468
```


Set power state (another capability) and end



Profit

- Analyzing PCIe is easy using radare2 - if you know what you're looking for
- When you design your PCIe driver - always check that its capabilities are working properly (look for DMA and unknown mappings)

```
0x0000444d    mov ecx, dword [n]
0x00004453    mov rdx, qword [local_938h]
0x0000445a    mov rsi, qword [local_930h]
0x00004461    mov rax, qword [local_928h]
0x00004468    mov r8d, 0
0x0000446e    mov rdi, rax
0x00004471    call sym.DeviceReadDMA
0x00004476    test eax, eax
0x00004478    jne 0x44a7
```

```
0x0000447a    mov ecx, dword [n]
0x00004480    mov rdx, qword [local_938h]
0x00004487    mov rsi, qword [local_930h]
0x0000448e    mov rax, qword [local_928h]
0x00004495    mov r8d, 0
0x0000449b    mov rdi, rax
0x0000449e    call sym.DeviceReadDMA
0x000044a3    test eax, eax
0x000044a5    je 0x44ae
```

Sum-up - Example

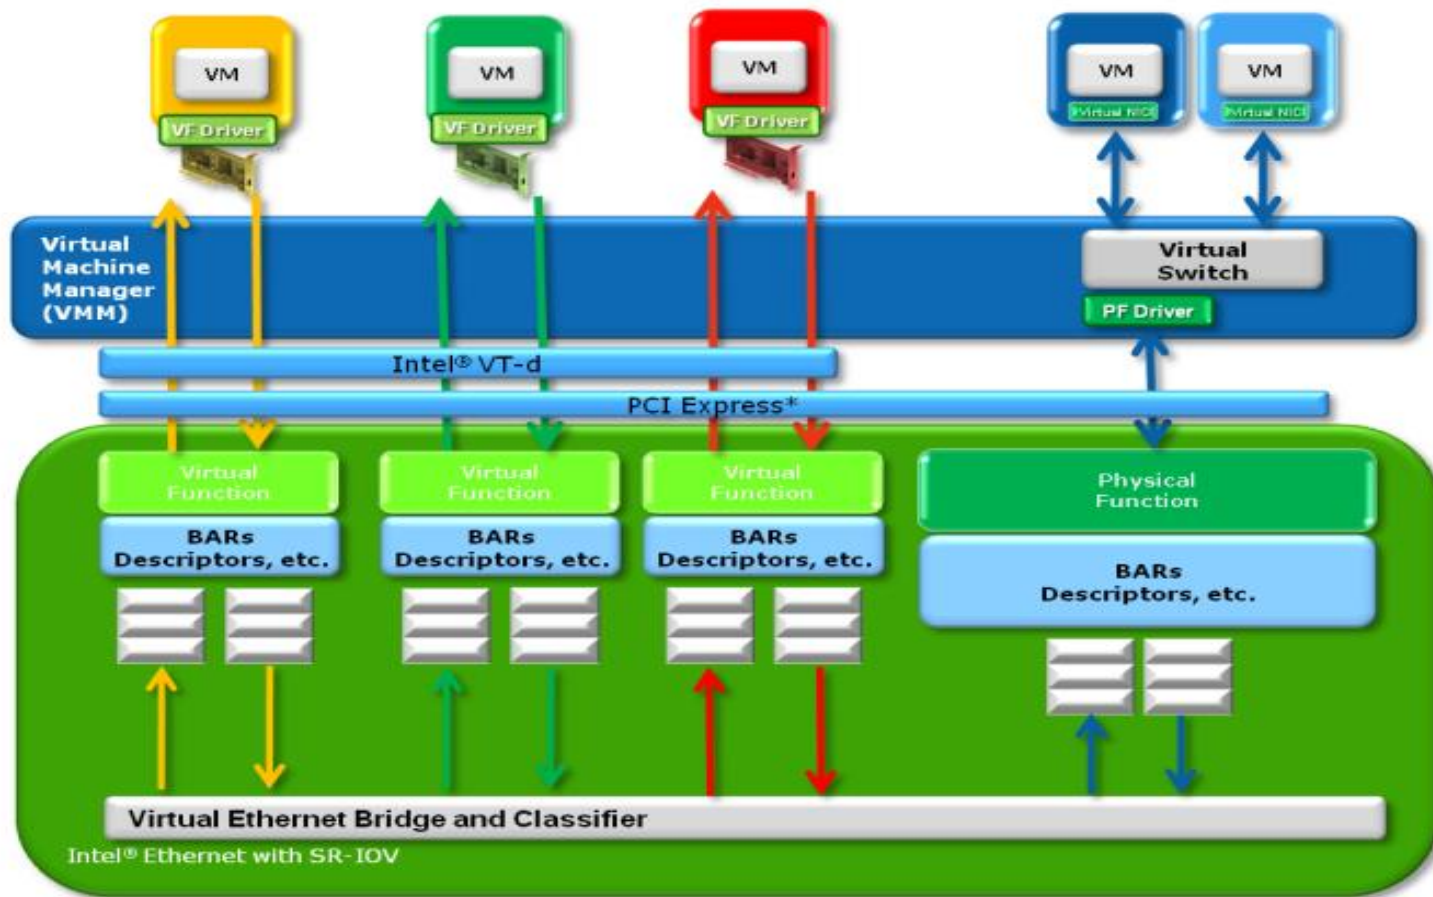
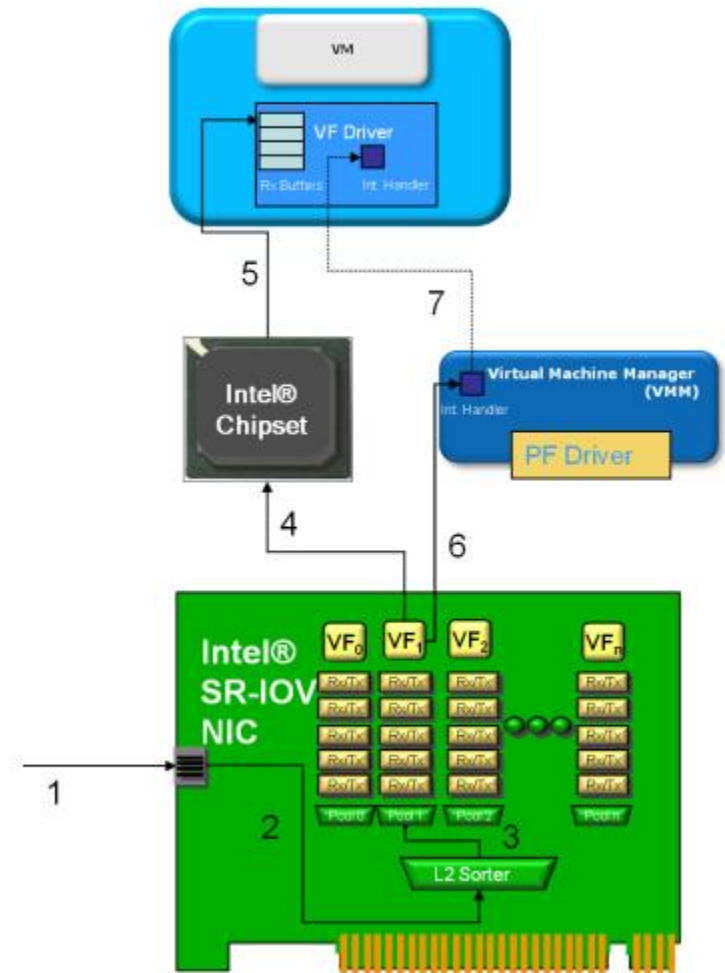


Figure from Inter PCI-SIG SR-IOV Primer


PCIe Data Path for incoming packets

1. The Ethernet packet arrives at the Ethernet NIC
2. The packet is sent to the Layer 2 sorter/switch/classifier
 - ▶ This Layer 2 sorter is configured by the Master Driver. When either the MD or the VF Driver configure a **MAC address** or **VLAN**, this Layer 2 sorter is configured.



Data Path for incoming packets

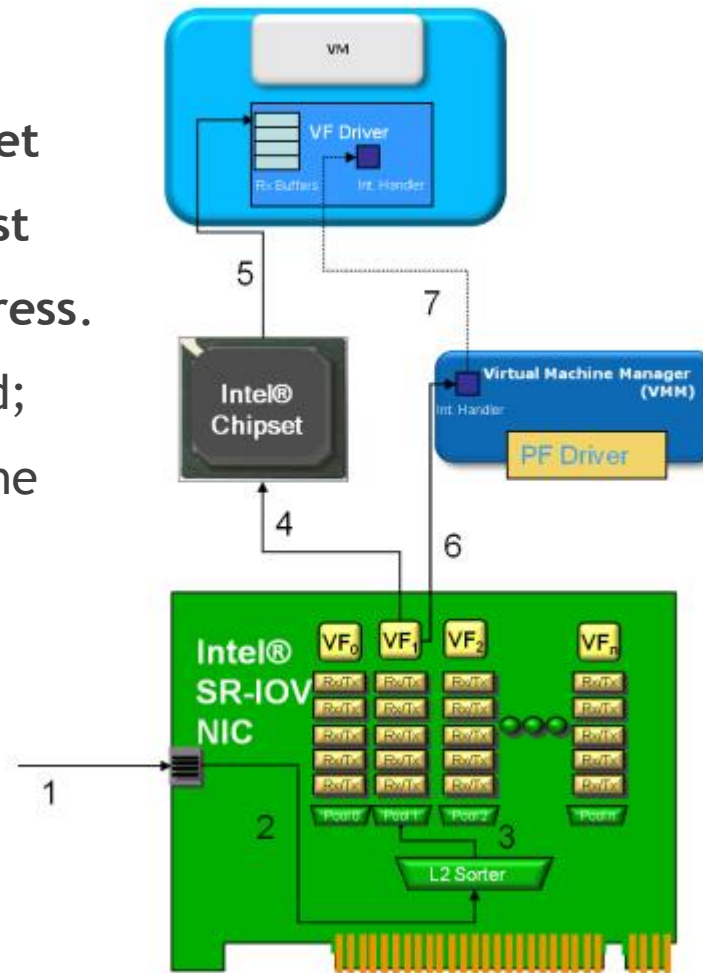
3. After being sorted by the Layer 2 Switch, the packet is placed into a receive queue dedicated to the target VF.
4. The DMA operation is initiated. The target memory address for the DMA operation is defined within the descriptors in the VF, which have been configured by the VF driver within the VM.



Data Path for incoming packets

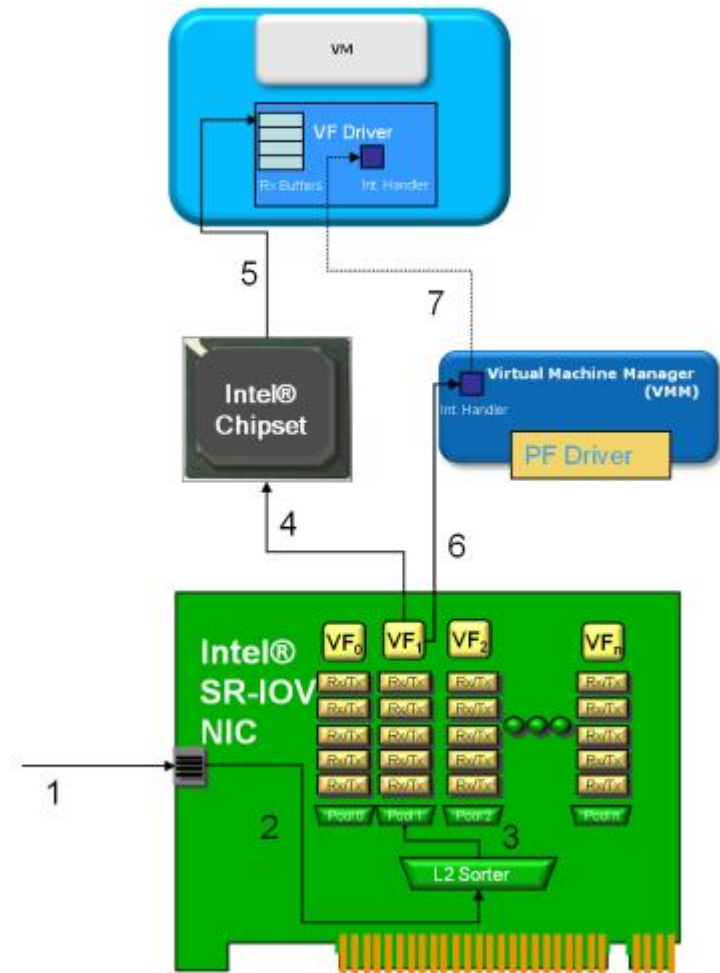
5. The DMA Operation has reached the chipset. Intel VT-d, which has been configured by the VMM then remaps the target DMA address from a virtual host address to a physical host address. The DMA operation is completed; the Ethernet packet is now in the memory space of the VM

6. The NIC fires interrupt, indicating a packet has arrived. This interrupt is handled by the VMM



Data Path for incoming packets

7. The VMM fires a virtual interrupt to the VM, so that it is informed that the packet has arrived



Now you can too!

- Find your favorite PCIe device
- Look for its device driver
- Check how the device is enumerated
- Look for its capabilities
- Follow the data flow and...
- Hack away!

Thanks!

- ▶ Adir Abraham
- ▶ Twitter: [@adirab](https://twitter.com/adirab)
- ▶ LinkedIn: <https://linkedin.com/in/adirab>

Reference

- PCI-SIG:

- PCI Express Base Specification v4.0 Revision 1.0
- Single Root I/O Virtualization and Sharing Specification Revision 1.1
- Address Translation Services Revision 1.1

- Intel PCI-SIG SR-IOV Primer

- <http://opensecuritytraining.info>

- VFIO: A User's perspective